

**Exam** : **Salesforce-AI-Specialist**

**Title** : Salesforce Certified AI  
Specialist Exam

**Vendor** : Salesforce

**Version** : V13.65

TOPEXAM.JP

**NO.1** Universal Containers (UC) wants to assess Salesforce's generative features but has concerns over its company data being exposed to third- party large language models (LLMs). Specifically, UC wants the following capabilities to be part of Einstein's generative AI service.

No data is used for LLM training or product improvements by third- party LLMs.

No data is retained outside of UC's Salesforce org.

The data sent cannot be accessed by the LLM provider.

Which property of the Einstein Trust Layer should the AI Specialist highlight to UC that addresses these requirements?

**A.** Prompt Defense

**B.** Zero-Data Retention Policy

**C.** Data Masking

**Answer:** B

Explanation:

Universal Containers (UC) has concerns about data privacy when using Salesforce's generative AI features, particularly around preventing third-party LLMs from accessing or retaining their data. The Zero- Data Retention Policy in the Einstein Trust Layer is designed to address these concerns by ensuring that:

- \* No data is used for training or product improvements by third-party LLMs.
- \* No data is retained outside of the customer's Salesforce organization.
- \* The LLM provider cannot access any customer data.

This policy aligns perfectly with UC's requirements for keeping their data safe while leveraging generative AI capabilities.

\* Prompt Defense and Data Masking are also security features, but they do not directly address the concerns related to third-party data access and retention.

References:

\* Salesforce Einstein Trust Layer Documentation: [https://help.salesforce.com/s/articleView?id=sf.einstein\\_trust\\_layer.htm](https://help.salesforce.com/s/articleView?id=sf.einstein_trust_layer.htm)

**NO.2** Universal Containers implemented Agentforce for its users. One user complains that an Agent is not deleting activities from the past 7 days. What is the reason for this issue?

**A.** Agentforce does not have the permission to delete the user's records.

**B.** Agentforce Delete Record Action permission is not associated to the user.

**C.** Agentforce does not have a standard Delete Record action.

**Answer:** C

Explanation:

\* Context of the Question Universal Containers (UC) uses Agentforce, a specialized AI-driven assistant for Salesforce. A user reports that an Agent is unable to delete recent activities.

\* Why Agentforce Cannot Delete Records

\* Agentforce's Standard Actions: Agentforce typically has predefined or "standard" actions like Create, Update, or Summarize records. However, a standard Delete Record action is not part of the default set of Agentforce actions.

\* Implication: If Agentforce has no built-in delete functionality, it cannot remove activities-even if the user has permission to delete them in the Salesforce UI.

\* Why Other Options Are Incorrect

\* Option A - Permission to Delete the User's Records: Standard Salesforce user permissions do not

automatically extend to Agentforce's capabilities. Even if the user can delete records, that doesn't grant Agentforce a new action.

\* Option B - Agentforce Delete Record Action Permission: There is no separate "Delete Record Action permission" for Agentforce to be toggled. The relevant issue is that the standard Delete Record action does not exist within Agentforce out of the box.

\* Conclusion The core reason for the issue is that Agentforce does not support a standard Delete Record action (Choice C).

Salesforce AI Specialist References & Documents

\* Salesforce Official Documentation - Agentforce (Note: Agentforce may be a pilot or specialized feature; check pilot release notes or official docs for standard actions.)

\* Salesforce AI Specialist Study Guide Covers the limitations of certain AI-enabled features regarding record operations.

**NO.3** Based on the user utterance, "Show me all the customers in New York", which standard Einstein Copilot action will the planner service use?

**A.** Query Records

**B.** Select Records

**C.** Fetch Records

**Answer: A**

Explanation:

The standard Einstein Copilot action that would be used in response to the user utterance, "Show me all the customers in New York," is Query Records. This action is responsible for retrieving a set of records from Salesforce based on a specified condition - in this case, filtering customers by location (New York).

\* Query Records is the action that fetches relevant data based on the criteria provided in the user's input.

\* Select Records is more about picking specific records from an already presented list.

\* Fetch Records is not a standard term used in this context for the action.

Refer to Einstein Copilot documentation on how Copilot actions work with natural language queries and data retrieval.

**NO.4** Universal Containers wants to use an external large language model (LLM) in Prompt Builder. What should an AI Specialist recommend?

**A.** Use Apex to connect to an external LLM and ground the prompt.

**B.** Use BYO-LLM functionality in Einstein Studio.

**C.** Use Flow and External Services to bring data from an external LLM.

**Answer: B**

Explanation:

Bring Your Own Large Language Model (BYO-LLM) functionality in Einstein Studio allows organizations to integrate and use external large language models (LLMs) within the Salesforce ecosystem.

Universal Containers can leverage this feature to connect and ground prompts with external LLMs, allowing for custom AI model use cases and seamless integration with Salesforce data.

\* Option B is the correct choice as Einstein Studio provides a built-in feature to work with external models.

- \* Option A suggests using Apex, but BYO-LLM functionality offers a more streamlined solution.
- \* Option C focuses on Flow and External Services, which is more about data integration and isn't ideal for working with LLMs.

References:

- \* Salesforce Einstein Studio BYO-LLM Documentation:

[https://help.salesforce.com/s/articleView?id=sf.einstein\\_studio\\_llm.htm](https://help.salesforce.com/s/articleView?id=sf.einstein_studio_llm.htm)

**NO.5** An AI Specialist is tasked to optimize a business process flow by assigning actions to agents within the Salesforce Agentforce Platform.

What is the correct method for the AI Specialist to assign actions to an Agent?

- A.** Assign the action to a Topic First in Agent Builder.
- B.** Assign the action to a Topic first on the Agent Actions detail page.
- C.** Assign the action to a Topic first on Action Builder.

**Answer:** C

- \* Action Builder is the central place in Salesforce Agentforce where you define and manage actions that your AI agents can perform. This includes connecting actions to various tools and systems.
- \* Topics in Agentforce represent the different tasks or intents that an AI agent can handle. By assigning an action to a Topic in Action Builder, you're essentially telling the agent, "When you encounter this type of request or situation, perform this action."

**NO.6** Universal Containers (UC) is Implementing Service AI Grounding to enhance its customer service operations.

UC wants to ensure that its AI-generated responses are grounded in the most relevant data sources. The team needs to configure the system to include all supported objects for grounding.

Which objects should UC select to configure Service AI Grounding?

- A.** Case, Knowledge, and Case Notes
- B.** Case and Knowledge
- C.** Case, Case Emails, and Knowledge

**Answer:** B

Explanation:

Universal Containers (UC) is implementing Service AI Grounding to enhance its customer service operations.

They aim to ensure that AI-generated responses are grounded in the most relevant data sources and need to configure the system to include all supported objects for grounding.

Supported Objects for Service AI Grounding:

- \* Case
- \* Knowledge
- \* Case Object:
  - \* Role in Grounding: Provides contextual data about customer inquiries, including case details, status, and history.
  - \* Benefit: Grounding AI responses in case data ensures that the information provided is relevant to the specific customer issue being addressed.
- \* Knowledge Object:
  - \* Role in Grounding: Contains articles and documentation that offer solutions and information

related to common issues.

- \* Benefit: Utilizing Knowledge articles helps the AI provide accurate and helpful responses based on verified information.
- \* Exclusion of Other Objects:
- \* Case Notes and Case Emails:
- \* Not Supported for Grounding: While useful for internal reference, these objects are not included in the supported objects for Service AI Grounding.
- \* Reason: They may contain sensitive or unstructured data that is not suitable for AI grounding purposes.

Why Options A and C are Incorrect:

- \* Option A (Case, Knowledge, and Case Notes):
- \* Case Notes Not Supported: Case Notes are not among the supported objects for grounding in Service AI.
- \* Option C (Case, Case Emails, and Knowledge):
- \* Case Emails Not Supported: Case Emails are also not included in the list of supported objects for grounding.

References:

- \* Salesforce AI Specialist Documentation - Service AI Grounding Configuration: Details the objects supported for grounding AI responses in Service Cloud.
- \* Salesforce Help - Implementing Service AI Grounding: Provides guidance on setting up grounding with Case and Knowledge objects.
- \* Salesforce Trailhead - Enhance Service with AI Grounding: Offers an interactive learning path on using AI grounding in service scenarios.

**NO.7** Which business requirement presents a good use case for leveraging Einstein Prompt Builder?

- A.** Forecast future sales trends based on historical data.
- B.** Identify potential high-value leads for targeted marketing campaigns.
- C.** Send reply to a request for proposal via a personalized email.

**Answer:** C

Explanation:

- \* Context of the Question
- \* Einstein Prompt Builder is a Salesforce feature that helps generate text (summaries, email content, responses) using AI models.
- \* The question presents three potential use cases, asking which one best fits the capabilities of Einstein Prompt Builder.
- \* Einstein Prompt Builder Typical Use Cases
- \* Text Generation & Summaries: Great for writing or summarizing content, like responding to an email or generating text for a record field.
- \* Why Not Forecast Future Sales Trends or Identify Potential High-Value Leads?
- \* (Option A) Forecasting trends typically involves predictive analytics and modeling capabilities found in Einstein Discovery or standard reporting, not generative text solutions.
- \* (Option B) Identifying leads for marketing campaigns involves lead scoring or analytics, again an Einstein Discovery or Lead Scoring scenario.
- \* Sending a Personalized RFP Email (Option C) is a classic example of using generative AI to compose well-structured, context-aware text.
- \* Conclusion Option C (Send reply to a request for proposal via a personalized email) is the best match

for Einstein Prompt Builder's generative text functionality.

Salesforce AI Specialist References & Documents

\* Salesforce Documentation: Einstein Prompt Builder Overview Highlights how to use Prompt Builder to create and customize text-based responses, especially for email or record fields.

\* Salesforce AI Specialist Study Guide Explains that generative AI features in Salesforce are designed for creating or summarizing text, not for advanced predictive use cases (like forecasting or lead scoring).

**NO.8** Universal Containers is evaluating Einstein Generative AI features to improve the productivity of the service center operation.

Which features should the AI Specialist recommend?

**A.** Service Replies and Case Summaries

**B.** Service Replies and Work Summaries

**C.** Reply Recommendations and Sales Summaries

**Answer: A**

Explanation:

To improve the productivity of the service center, the AI Specialist should recommend the Service Replies and Case Summaries features.

\* Service Replies helps agents by automatically generating suggested responses to customer inquiries, reducing response time and improving efficiency.

\* Case Summaries provide a quick overview of case details, allowing agents to get up to speed faster on customer issues.

\* Work Summaries are not as relevant for direct customer service operations, and Sales Summaries are focused on sales processes, not service center productivity.

For more information, see Salesforce's Einstein Service Cloud documentation on the use of generative AI to assist customer service teams.

**NO.9** Which configuration must an AI Specialist complete for users to access generative AI-enabled fields in the Salesforce mobile app?

**A.** Enable Mobile Generative AI.

**B.** Enable Mobile Prompt Responses.

**C.** Enable Dynamic Forms on Mobile.

**Answer: A**

Explanation:

\* Context of the Question

\* Universal Containers (UC) has generative AI-enabled fields that users can access in the desktop experience.

\* The AI Specialist needs these same fields to be visible and usable in the Salesforce Mobile App.

\* Why Dynamic Forms on Mobile?

\* Dynamic Forms allow you to configure record pages so that fields and sections can appear or be hidden based on certain criteria.

\* When you enable "Dynamic Forms for Mobile," any generative AI-enabled fields placed on the dynamic layout become accessible in the Salesforce mobile experience.

\* There is no standard Setup option labeled "Enable Mobile Generative AI" or "Enable Mobile Prompt Responses" as a universal toggle; the existing official approach is to ensure dynamic forms (and the

relevant fields) are supported on mobile.

\* Conclusion

\* Ensuring that these AI-driven fields are visible on mobile is accomplished by turning on Dynamic Forms on Mobile and adding those fields to the dynamic layout. Therefore, Option C is correct.

Salesforce AI Specialist References & Documents

\* Salesforce Documentation: Dynamic Forms Overview Explains how to enable Dynamic Forms for both desktop and mobile UIs, allowing newly added fields (including generative AI-enabled ones) to display in the Salesforce Mobile App.

\* Salesforce AI Specialist Study Guide Reiterates that to expose generative AI fields or components in mobile, you must configure dynamic forms and ensure compatibility on mobile layouts.

**NO.10** Universal Containers (UC) wants to create a new Sales Email prompt template in Prompt Builder using the

"Save As" function. However, UC notices that the new template produces different results compared to the standard Sales Email prompt due to missing hyperparameters.

What should UC do to ensure the new prompt template produces results comparable to the standard Sales Email prompts?

**A.** Use Model Playground to create a model configuration with the specified parameters.

**B.** Manually add the hyperparameters to the new template.

**C.** Revert to using the standard template without modifications.

**Answer:** B

Explanation:

When Universal Containers creates a new Sales Email prompt template using the "Save As" function, missing hyperparameters can result in different outputs. To ensure the new prompt produces comparable results to the standard Sales Email prompt, the AI Specialist should manually add the necessary hyperparameters to the new template.

\* Hyperparameters like Temperature, Frequency Penalty, and Presence Penalty directly affect how the AI generates responses. Ensuring that these are consistent with the standard template will result in similar outputs.

\* Option A (Model Playground) is not necessary here, as it focuses on fine-tuning models, not adjusting templates directly.

\* Option C (Reverting to the standard template) does not solve the issue of customizing the prompt template.

For more information, refer to Prompt Builder documentation on configuring hyperparameters in custom templates.

**NO.11** Universal Containers (UC) has a mature Salesforce org with a lot of data in cases and Knowledge articles. UC is concerned that there are many legacy fields, with data that might not be applicable for Einstein AI to draft accurate email responses.

Which solution should UC use to ensure Einstein AI can draft responses from a defined data source?

**A.** Service AI Grounding

**B.** Work Summaries

**C.** Service Replies

**Answer:** A

Explanation:

Service AI Grounding is the solution that Universal Containers should use to ensure Einstein AI drafts responses based on a well-defined data source. Service AI Grounding allows the AI model to be anchored in specific, relevant data sources, ensuring that any AI-generated responses (e.g., email replies) are accurate, relevant, and drawn from up-to-date information, such as Knowledge articles or cases.

Given that UC has legacy fields and outdated data, Service AI Grounding ensures that only the valid and applicable data is used by Einstein AI to craft responses. This helps improve the relevance of responses and avoids inaccuracies caused by outdated or irrelevant fields.

Work Summaries and Service Replies are useful features but do not address the need for grounding AI outputs in specific, current data sources like Service AI Grounding does.

For more details, you can refer to Salesforce's Service AI Grounding documentation for managing AI-generated content based on accurate data sources.

**NO.12** Universal Containers (UC) wants to improve the productivity of its sales team with generative AI technology.

However, UC is concerned that public AI virtual assistants lack adequate company data to general useful responses.

Which solution should UC consider?

- A.** fine-tune the Einstein AI model with CBM data.
- B.** Build AI model with Einstein discovery and deploy to sales users.
- C.** Enable Agentforce and deploy to sales users.

**Answer:** A

Explanation:

\* Context of the Question: Universal Containers (UC) wants to harness generative AI to boost sales productivity. They are wary of public AI virtual assistants (like generic chatbots) that lack sufficient UC-specific data to generate useful business responses.

\* Why Fine-Tune an Einstein AI Model with CRM Data?

\* Company-Specific Relevance: By fine-tuning Einstein AI with UC's CRM data (accounts, opportunities, products, and historical interactions), the model learns the enterprise-specific context. This ensures that the generative outputs are accurate and tailored to UC's sales scenarios.

\* Security and Compliance: Using Salesforce Einstein within the Salesforce ecosystem keeps data under UC's control, aligning with trust, security, and compliance requirements.

\* Better Predictions: Einstein AI can produce more relevant insights (e.g., recommended next steps, content suggestions, or AI-generated email responses) when it has been trained on real, high-quality internal data.

\* Why Not Build an AI Model with Einstein Discovery (Option B)?

\* Einstein Discovery Use Case: Einstein Discovery is best suited for predictive and prescriptive analytics (e.g., analyzing large data sets for patterns, scoring leads, or predicting churn). While it provides advanced analytics, it is not primarily designed for generative text-based interactions for end-user consumption in a conversational format.

\* Why Not Enable Agentforce (Option C)?

\* Agentforce Overview: "Agentforce" (sometimes referencing a pilot or non-mainstream name) typically focuses on interactive help or workforce collaboration. It does not inherently solve the problem of large-scale generative AI using internal CRM data. Moreover, you still need a robust generative engine fine-tuned on company data.

\* Outcome: Fine-tuning the Einstein AI model with UC's CRM data (Answer A) is the most direct,



Salesforce-native solution to provide generative AI responses that are aligned with UC's context, driving productivity gains and ensuring data privacy.

Salesforce AI Specialist References & Documents

- \* Salesforce Official: Einstein GPT Overview
- \* Discusses how Einstein GPT can be fine-tuned with specific CRM data to deliver contextually relevant, generative AI responses.
- \* Salesforce Trailhead: Get Started with Salesforce Einstein
- \* Explains the fundamentals of AI within the Salesforce platform, including training and optimizing Einstein models.
- \* Salesforce Documentation: Einstein Discovery
- \* Details how Einstein Discovery is primarily used for advanced analytics and predictions, not direct generative text solutions.
- \* Salesforce AI Specialist Study Guide
- \* Provides the official outline of Einstein AI capabilities, referencing how to configure and fine-tune models for specialized enterprise use cases.

**NO.13** Universal Containers (UC) wants to use Flow to bring data from unified Data Cloud objects to prompt templates.

Which type of flow should UC use?

- A.** Data Cloud-triggered flow
- B.** Template-triggered prompt flow
- C.** Unified-object linking flow

**Answer:** B

Explanation:

In this scenario, Universal Containers wants to bring data from unified Data Cloud objects into prompt templates, and the best way to do that is through a Data Cloud-triggered flow. This type of flow is specifically designed to trigger actions based on data changes within Salesforce Data Cloud objects.

Data Cloud-triggered flows can listen for changes in the unified data model and automatically bring relevant data into the system, making it available for prompt templates. This ensures that the data is both real-time and up-to-date when used in generative AI contexts.

For more detailed guidance, refer to Salesforce documentation on Data Cloud-triggered flows and Data Cloud integrations with generative AI solutions.

**NO.14** An AI Specialist wants to use the related lists from an account in a custom prompt template.

What should the AI Specialist consider when configuring the prompt template?

- A.** The text encoding (for example, UTF-8, ASCII) option
- B.** The maximum number of related list merge fields
- C.** The choice between XML and JSON rendering formats for the list

**Answer:** B

Explanation:

When configuring a custom prompt template to use related lists, the AI Specialist must be aware of the maximum number of related list merge fields that can be included. Salesforce enforces limits to ensure prompt templates perform efficiently and do not overload the system with too much data. As a best practice, it's important to monitor and optimize the number of merge fields used.

\* Option B is correct because there is a limit on how many related list merge fields can be included in a prompt template.

\* Option A (text encoding) and Option C (XML/JSON rendering) are not key considerations in this context.

References:

\* Salesforce Prompt Builder Documentation: [https://help.salesforce.com/s/articleView?id=sf.prompt\\_builder.htm](https://help.salesforce.com/s/articleView?id=sf.prompt_builder.htm)

**NO.15** Universal Containers tests out a new Einstein Generative AI feature for its sales team to create personalized and contextualized emails for its customers. Sometimes, users find that the draft email contains placeholders for attributes that could have been derived from the recipient's contact record.

What is the most likely explanation for why the draft email shows these placeholders?

**A.** The user does not have Einstein Sales Emails permission assigned.

**B.** The user does not have permission to access the fields.

**C.** The user's locale language is not supported by Prompt Builder.

**Answer: B**

Explanation:

When using Einstein Generative AI to create personalized emails, if placeholders appear in the draft email where data from a recipient's Contact record should be, the most likely reason is that the user lacks permission to access the necessary fields. Salesforce's field-level security may prevent users from viewing or utilizing certain data fields, resulting in placeholders being shown instead of the actual values.

\* Option B is correct because missing field permissions will cause placeholders in email drafts.

\* Option A (missing Einstein Sales Emails permission) is unlikely, as this would prevent email generation altogether, not just placeholders.

\* Option C (locale language issues) would more likely affect language-specific issues, not field placeholders.

References:

\* Salesforce Email Template and Permissions Documentation:

[https://help.salesforce.com/s/articleView?](https://help.salesforce.com/s/articleView?id=sf.email_templates_field_permissions.htm)

[id=sf.email\\_templates\\_field\\_permissions.htm](https://help.salesforce.com/s/articleView?id=sf.email_templates_field_permissions.htm)

**NO.16** Universal Containers (UC) uses Salesforce Service Cloud to support its customers and agents handling cases.

UC is considering implementing Einstein Copilot and extending Service Cloud to mobile users.

When would Einstein Copilot implementation be most advantageous?

**A.** When the goal is to streamline customer support processes and improve response times

**B.** When the main objective is to enhance data security and compliance measures

**C.** When the focus is on optimizing marketing campaigns and strategies

**Answer: A**

Explanation:

Einstein Copilot implementation would be most advantageous in Salesforce Service Cloud when the goal is to streamline customer support processes and improve response times. Einstein Copilot can assist agents by providing real-time suggestions, automating repetitive tasks, and generating

contextual responses, thus enhancing service efficiency.

\* Option B (data security) is not the primary focus of Einstein Copilot, which is more about improving operational efficiency.

\* Option C (marketing campaigns) falls outside the scope of Service Cloud and Einstein Copilot's primary benefits, which are aimed at improving customer service and case management.

For further reading, refer to Salesforce documentation on Einstein Copilot for Service Cloud and how it improves support processes.

**NO.17** Universal Containers wants to allow its service agents to query the current fulfillment status of an order with natural language. There is an existing auto launched flow to query the information from Oracle ERP, which is the system of record for the order fulfillment process.

How should an AI Specialist apply the power of conversational AI to this use case?

**A.** Create a Flex prompt template in Prompt Builder.

**B.** Create a custom copilot action which calls a flow.

**C.** Configure the Integration Flow Standard Action in Einstein Copilot.

**Answer:** B

Explanation:

To enable Universal Containers service agents to query the current fulfillment status of an order using natural language and leverage an existing auto-launched flow that queries Oracle ERP, the best solution is to create a custom copilot action that calls the flow. This action will allow Einstein Copilot to interact with the flow and retrieve the required order fulfillment information seamlessly. Custom copilot actions can be tailored to call various backend systems or flows in response to user requests.

\* Option B is correct because it enables integration between Einstein Copilot and the flow that connects to Oracle ERP.

\* Option A (Flex prompt template) is more suited for static responses and not for invoking flows.

\* Option C (Integration Flow Standard Action) is not directly related to creating a specific copilot action for this use case.

References:

\* Salesforce Einstein Copilot Actions: [https://help.salesforce.com/s/articleView?id=einstein\\_copilot\\_actions.htm](https://help.salesforce.com/s/articleView?id=einstein_copilot_actions.htm)

**NO.18** An AI Specialist is tasked with analyzing Agent interactions looking into user inputs, requests, and queries to identify patterns and trends.

What functionality allows the AI Specialist to achieve this?

**A.** User Utterances dashboard

**B.** Agent Event Logs dashboard

**C.** AI Audit & Feedback Data dashboard

**Answer:** A

Explanation:

The User Utterances dashboard (Option A) is the correct functionality for analyzing user inputs, requests, and queries to identify patterns and trends. This dashboard aggregates and categorizes the natural language inputs (utterances) from users, enabling the AI Specialist to:

\* Identify Common Queries: Surface frequently asked questions or recurring issues.

\* Detect Intent Patterns: Understand how users phrase requests, which helps refine intent detection models.

\* Improve Bot Training: Highlight gaps in training data or misclassified utterances that require adjustment.

Why Other Options Are Incorrect:

\* B. Agent Event Logs dashboard: Focuses on agent activity (e.g., response times, resolved cases) rather than user input analysis.

\* C. AI Audit & Feedback Data dashboard: Tracks AI model performance, audit trails, and user feedback scores but does not directly analyze raw user utterances or queries.

References:

\* Salesforce Einstein AI Specialist Certification Guide: Emphasizes the User Utterances dashboard as the primary tool for analyzing user inputs to improve conversational AI.

\* Trailhead Module: "Einstein Bots Basics" highlights using the dashboard to refine bot training based on user interaction data.

\* Salesforce Help Documentation: Describes the User Utterances dashboard as critical for identifying trends in customer interactions.

**NO.19** The sales team at a hotel resort would like to generate a guest summary about the guests' interests and provide recommendations based on their activity preferences captured in each guest profile. They want the summary to be available only on the contact record page.

Which AI capability should the team use?

**A.** Einstein Copilot

**B.** Prompt Builder

**C.** Model Builder

**Answer:** B

Explanation:

The sales team at a hotel resort wants to generate a guest summary about guests' interests and provide recommendations based on their activity preferences captured in each guest profile. They require the summary to be available only on the contact record page.

Solution:

\* Use Prompt Builder to create a prompt template that generates the desired summary and displays it on the contact record page.

\* Prompt Builder:

\* Purpose: Allows the creation of custom prompt templates that leverage AI to generate content based on Salesforce data.

\* Functionality:

\* Field Generation Templates: Can be used to populate fields on records with AI-generated summaries.

\* Customization: Enables the AI Specialist to design prompts that utilize data from the guest profiles to produce personalized summaries and recommendations.

\* Relevance to the Use Case:

\* The sales team wants the summary to be available on the contact record page, which aligns with the capabilities of Prompt Builder to generate and display content on specific record pages.

\* Implementation Steps:

\* Create a Field Generation Prompt Template:

\* Use Prompt Builder to create a new prompt template of type Field Generation.

\* Design the prompt to instruct the AI to generate a summary based on the guest's interests and activity preferences.

- \* Include Relevant Data:
- \* Use merge fields to include data from the guest profile in the prompt.
- \* Ensure that the prompt accesses the necessary fields to generate accurate recommendations.
- \* Configure the Contact Page Layout:
- \* Add the field that will display the AI-generated summary to the contact record page layout.
- \* Ensure that the field is only visible where appropriate, adhering to the requirement of availability only on the contact record page.
- \* Why Not Einstein Copilot or Model Builder:
- \* Option A (Einstein Copilot):
- \* Purpose: Einstein Copilot is a conversational AI assistant designed to interact with users through natural language.
- \* Mismatch with Requirements:
- \* The team wants a static summary displayed on the contact record page, not an interactive conversational experience.
- \* Option C (Model Builder):
- \* Purpose: Model Builder is used to create custom AI models for predictions and classifications.
- \* Inapplicability:
- \* Building a custom model is unnecessary for generating text summaries based on existing data.
- \* Model Builder does not directly provide functionality to generate and display summaries on record pages.

#### References:

- \* Salesforce AI Specialist Documentation - Prompt Builder Overview:
- \* Provides an introduction to Prompt Builder and its capabilities.
- \* Salesforce Help - Creating Field Generation Prompt Templates:
- \* Guides on creating prompt templates that generate content for fields on records.
- \* Salesforce Trailhead - Customize AI Content with Prompt Builder:
- \* Offers hands-on experience in building and customizing prompt templates.

#### Conclusion:

By utilizing Prompt Builder, the sales team can create a customized prompt template that generates personalized guest summaries and recommendations based on activity preferences. This solution meets the requirement of displaying the summary only on the contact record page, enhancing the team's ability to engage with guests effectively.

**NO.20** Universal Containers (UC) is experimenting with using public Generative AI models and is familiar with the language required to get the information it needs. However, it can be time consuming for both UC's sales and service reps to type in the prompt to get the information they need, and ensure prompt consistency.

Which Salesforce feature should a Salesforce AI Specialist recommend to address these concerns?

- A.** Einstein Recommendation Builder
- B.** Einstein Copilot Action: Query Records
- C.** Einstein Prompt Builder and Prompt Templates

**Answer:** C

#### Explanation:

For Universal Containers (UC), to reduce the time and ensure prompt consistency when using public generative AI models, the recommended feature is Einstein Prompt Builder and Prompt Templates. This feature allows teams to create reusable and consistent prompts for generative AI tasks, ensuring

that all users receive uniform responses without having to type in detailed prompts manually every time.

\* Einstein Prompt Builder simplifies the creation of prompts, and Prompt Templates standardize the inputs, saving time for sales and service reps.

\* Option A (Einstein Recommendation Builder) is more focused on recommendations, not prompt standardization.

\* Option B (Einstein Copilot Action: Query Records) is for querying records, not generating AI-driven prompts.

References:

\* Salesforce Prompt Builder Overview: [https://help.salesforce.com/s/articleView?id=sf.prompt\\_builder\\_overview.htm](https://help.salesforce.com/s/articleView?id=sf.prompt_builder_overview.htm)

**NO.21** An AI Specialist configured Data Masking within the Einstein Trust Layer.

How should the AI Specialist begin validating that the correct fields are being masked?

**A.** Use a Flow-based resource in Prompt Builder to debug the fields' merge values using Flow Debugger.

**B.** Request the Einstein Generative AI Audit Data from the Security section of the Setup menu.

**C.** Enable the collection and storage of Einstein Generative AI Audit Data on the Einstein Feedback setup page.

**Answer:** C

Explanation:

To begin validating that the correct fields are being masked in Einstein Trust Layer, the AI Specialist should request the Einstein Generative AI Audit Data from the Security section of the Salesforce Setup menu.

This audit data allows the AI Specialist to see how data is being processed, including which fields are being masked, providing transparency and validation that the configuration is working as expected.

\* Option B is correct because it allows for the retrieval of audit data that can be used to validate data masking.

\* Option A (Flow Debugger) and Option C (Einstein Feedback) do not relate to validating field masking in the context of the Einstein Trust Layer.

References:

\* Salesforce Einstein Trust Layer Documentation: [https://help.salesforce.com/s/articleView?id=sf.einstein\\_trust\\_layer\\_audit.htm](https://help.salesforce.com/s/articleView?id=sf.einstein_trust_layer_audit.htm)

**NO.22** How does Secure Data Retrieval ensure that only authorized users can access necessary Salesforce data for dynamic grounding?

**A.** Retrieves Salesforce data based on the 'Run As' users permissions.

**B.** Retrieves Salesforce data based on the user's permissions executing the prompt.

**C.** Retrieves Salesforce data based on the Prompt template's object permissions.

**Answer:** B

Explanation:

Secure Data Retrieval enforces Salesforce's security model by dynamically grounding data access in the permissions of the user executing the prompt. This ensures compliance with CRUD (Create, Read, Update, Delete) and FLS (Field-Level Security) settings, preventing unauthorized access to sensitive data. For example, if a user lacks access to a specific object or field, the AI model cannot retrieve it

for dynamic grounding.

- \* "Run As" user permissions (A) would bypass user-specific security, posing a compliance risk.
- \* Prompt template permissions (C) are not a Salesforce security mechanism; access is always tied to the user's profile and sharing settings.

**NO.23** A sales manager is using Agent Assistant to streamline their daily tasks. They ask the agent to Show me a list of my open opportunities.

How does the large language model (LLM) in Agentforce identify and execute the action to show the sales manager a list of open opportunities?

- A.** The LLM interprets the user's request, generates a plan by identifying the appropriate topics and actions, and executes the actions to retrieve and display the open opportunities
- B.** The LLM uses a static set of rules to match the user's request with predefined topics and actions, bypassing the need for dynamic interpretation and planning.
- C.** Using a dialog pattern, the LLM matches the user query to the available topic, action and steps then performs the steps for each action, such as retrieving a list of open opportunities.

**Answer: A**

Explanation:

Agentforce's LLM dynamically interprets natural language requests (e.g., "Show me open opportunities"), generates an execution plan using the planner service, and retrieves data via actions (e.g., querying Salesforce records). This contrasts with static rules (B) or rigid dialog patterns (C), which lack contextual adaptability. Salesforce documentation highlights the planner's role in converting intents into actionable steps while adhering to security and business logic.

**NO.24** An AI Specialist has grounded a prompt template with a related list. During user acceptance testing (UAT).

users are not getting the correct responses.

What is causing this issue?

- A.** The related list is not on the parent object's page layout.
- B.** The related list is Read Only.
- C.** The related list prompt template option is not enabled.

**Answer: C**

Explanation:

When grounding a prompt template with a related list, the AI must be explicitly configured to include the related list's data. If the "related list prompt template option" is not enabled, the AI ignores the related list, leading to incomplete or incorrect responses.

- \* Option A: Page layout visibility affects user interface display but does not restrict data access for AI grounding.
- \* Option B: Read-only settings prevent edits but not data retrieval.
- \* Option C: Enabling the related list in the prompt template configuration is mandatory for the AI to use its data.

References:

- \* Salesforce Help: Prompt Template Grounding Settings
- \* States that "related lists must be enabled in the prompt template's grounding settings to include their data in AI responses."

**NO.25** Universal Containers' data science team is hosting a generative large language model (LLM) on Amazon Web Services (AWS).

What should the team use to access externally-hosted models in the Salesforce Platform?

- A. Model Builder
- B. App Builder
- C. Copilot Builder

**Answer:** A

Explanation:

To access externally-hosted models, such as a large language model (LLM) hosted on AWS, the Model Builder in Salesforce is the appropriate tool. Model Builder allows teams to integrate and deploy external AI models into the Salesforce platform, making it possible to leverage models hosted outside of Salesforce infrastructure while still benefiting from the platform's native AI capabilities.

\* Option B, App Builder, is primarily used to build and configure applications in Salesforce, not to integrate AI models.

\* Option C, Copilot Builder, focuses on building assistant-like tools rather than integrating external AI models.

Model Builder enables seamless integration with external systems and models, allowing Salesforce users to use external LLMs for generating AI-driven insights and automation.

Salesforce AI Specialist References: For more details, check the Model Builder guide here:

[https://help.](https://help.salesforce.com/s/articleView?id=sf.model_builder_external_models.htm)

[salesforce.com/s/articleView?id=sf.model\\_builder\\_external\\_models.htm](https://help.salesforce.com/s/articleView?id=sf.model_builder_external_models.htm)

**NO.26** Universal Containers (UC) is building a Flex prompt template. UC needs to use data returned by the flow in the prompt template.

Which flow element should UC use?

- A. Add Flow Instructions
- B. Add Prompt Instructions
- C. Add Flex Instructions

**Answer:** A

**NO.27** An AI Specialist wants to include data from the response of external service invocation (REST API callout) into the prompt template.

How should the AI Specialist meet this requirement?

- A. Convert the JSON to an XML merge field.
- B. Use External Service Record merge fields.
- C. Use "Add Prompt Instructions" flow element.

**Answer:** B

Explanation:

An AI Specialist wants to include data from the response of an external service invocation (REST API callout) into a prompt template. The goal is to incorporate dynamic data retrieved from an external API into the AI-generated content.

Solution:

\* Use External Service Record Merge Fields

\* External Service Integration:

\* Definition: External Services in Salesforce allow the integration of external REST APIs into Salesforce



without custom code.

- \* Registration: The external service must be registered in Salesforce, defining the API's schema and methods.

- \* External Service Record Merge Fields:

- \* Purpose: Enables the inclusion of data from external service responses directly into prompt templates using merge fields.

- \* Functionality:

- \* Dynamic Data Inclusion: Allows prompt templates to access and use data returned from REST API callouts.

- \* Merge Fields Syntax: Use merge fields in the prompt template to reference specific data points from the API response.

Implementation Steps:

- \* Register the External Service:

- \* Use External Services to register the REST API in Salesforce.

- \* Define the API's schema, including methods and data structures.

- \* Create a Named Credential:

- \* Configure authentication and endpoint details for the external API.

- \* Use External Service in Flow:

- \* Build a Flow that invokes the external service and captures the response.

- \* Ensure the flow outputs the necessary data for use in the prompt template.

- \* Configure the Prompt Template:

- \* Use External Service Record merge fields in the prompt template to reference data from the flow's output.

- \* Syntax Example: {{flowOutputVariable.fieldName}}

Why Other Options are Less Suitable:

- \* Option A (Convert the JSON to an XML merge field):

- \* Irrelevance: Converting JSON to XML merge fields is unnecessary and complicates the process.

- \* Unsupported Method: Salesforce prompt templates do not support direct inclusion of XML merge fields from JSON conversion.

- \* Option C (Use "Add Prompt Instructions" flow element):

- \* Purpose of Add Prompt Instructions:

- \* Allows adding instructions to the prompt within a flow but does not facilitate including external data.

- \* Limitation: Does not directly help in incorporating external service responses into the prompt template.

References:

- \* Salesforce AI Specialist Documentation - Integrating External Services with Prompt Templates:

- \* Explains how to use External Services and merge fields in prompt templates.

- \* Salesforce Help - Using Merge Fields with External Data:

- \* Provides guidance on referencing external data in templates using merge fields.

- \* Salesforce Trailhead - External Services and Flow:

- \* Offers a practical understanding of integrating external APIs using External Services and Flow.

Conclusion:

By using External Service Record merge fields, the AI Specialist can effectively include data from external REST API responses into prompt templates, ensuring that the AI-generated content is enriched with up-to-date and relevant external data.

**NO.28** Universal Containers plans to enhance its sales team's productivity using AI. Which specific requirement necessitates the use of Prompt Builder?

- A.** Predicting the likelihood of customers churning or discontinuing their relationship with the company
- B.** Creating a draft newsletter for an upcoming tradeshow
- C.** Creating an estimated Customer Lifetime Value (CLV) with historical purchase data

**Answer:** B

**NO.29** When a customer chat is initiated, which functionality in Salesforce provides generative AI replies or draft emails based on recommended Knowledge articles?

- A.** Einstein Reply Recommendations
- B.** Einstein Service Replies
- C.** Einstein Grounding

**Answer:** B

Explanation:

When a customer chat is initiated, Einstein Service Replies provides generative AI replies or draft emails based on recommended Knowledge articles. This feature uses the information from the Salesforce Knowledge base to generate responses that are relevant to the customer's query, improving the efficiency and accuracy of customer support interactions.

\* Option B is correct because Einstein Service Replies is responsible for generating AI-driven responses based on knowledge articles.

\* Option A (Einstein Reply Recommendations) is focused on recommending replies but does not generate them.

\* Option C (Einstein Grounding) refers to grounding responses in data but is not directly related to drafting replies.

References:

\* Einstein Service Replies Overview: [https://help.salesforce.com/s/articleView?id=sf.einstein\\_service\\_replies.htm](https://help.salesforce.com/s/articleView?id=sf.einstein_service_replies.htm)

**NO.30** Universal Containers has an active standard email prompt template that does not fully deliver on the business requirements.

Which steps should an AI Specialist take to use the content of the standard prompt email template in question and customize it to fully meet the business requirements?

- A.** Save as New Template and edit as needed.
- B.** Clone the existing template and modify as needed.
- C.** Save as New Version and edit as needed.

**Answer:** A

Explanation:

When an active standard email prompt template doesn't meet the business requirements, the best approach is to clone the existing template and modify it as needed. Cloning allows the AI Specialist to preserve the original template while making adjustments to fit specific business needs. This ensures that any customizations are applied without altering the original standard template.

Saving as a new version is typically used for versioning changes in the same template, while Save as New Template creates a brand-new template without linking to the existing one. Cloning provides a

balance, allowing modifications while retaining the original structure for future reference. For more details, refer to Salesforce Prompt Builder documentation for guidance on cloning and modifying templates.

**NO.31** Universal Containers wants support agents to use Agentforce to ask questions about its product tutorials and product guides.

What should the AI Specialist do to meet this requirement?

- A.** Create a prompt template for product tutorials and guides.
- B.** Add an Answer Questions custom field in the product object for tutorial instructions.
- C.** Publish product tutorials and guides as Knowledge articles.

**Answer:** C

Explanation:

\* Context of the Question Universal Containers (UC) wants its support agents to use Agentforce to ask questions about product tutorials and product guides. Agentforce typically references knowledge sources to provide accurate and contextual responses.

\* Why Knowledge Articles?

\* Centralized Repository: Publishing product tutorials and guides as Knowledge articles in Salesforce ensures that the information is readily available and searchable by Agentforce.

\* AI Integration: Salesforce's AI solutions, including Agentforce, can often be configured to pull content directly from Salesforce Knowledge articles, giving users on-demand answers without manual data duplication.

\* Maintenance & Updates: Storing content in Salesforce Knowledge simplifies content updates, versioning, and user permissions.

\* Why Not the Other Options?

\* Option A (Create a Prompt Template): Creating a prompt template alone does not solve how the underlying content (tutorials, guides) is stored or accessed by Agentforce. Prompt templates shape the queries/responses but do not provide the knowledge base.

\* Option B (Add an Answer Questions Custom Field): A single field on the product object is insufficient for the depth of information found in tutorials and guides. It also lacks the robust search and user-friendly interface that Knowledge articles provide.

\* Conclusion To ensure Agentforce can effectively retrieve and deliver accurate information about products, publishing product tutorials and guides as Knowledge articles is the recommended approach.

Salesforce AI Specialist References & Documents

\* Salesforce Documentation: Set Up Salesforce Knowledge Discusses how to publish articles for easy access by AI-driven assistants and support teams.

\* Salesforce AI Specialist Study Guide Explains best practices for feeding knowledge sources to generative AI and Agentforce.

**NO.32** What is the correct process to leverage Prompt Builder in a Salesforce org?

- A.** Select the appropriate prompt template type to use, select one of Salesforce's standard prompts, determine the object to associate the prompt, select a record to validate against, and associate the prompt to an action.
- B.** Select the appropriate prompt template type to use, develop the prompt within the prompt workspace, select resources to dynamically insert CRM-derived grounding data, pick the model to use, and test and validate the generated responses.

**C.** Enable the target object for generative prompting, develop the prompt within the prompt workspace, select records to fine-tune and ground the response, enable the Trust Layer, and associate the prompt to an action.

**Answer:** B

Explanation:

When using Prompt Builder in a Salesforce org, the correct process involves several important steps:

- \* Select the appropriate prompt template type based on the use case.
- \* Develop the prompt within the prompt workspace, where the template is created and customized.
- \* Select CRM-derived grounding data to be dynamically inserted into the prompt, ensuring that the AI-generated responses are based on accurate and relevant data.
- \* Pick the model to use for generating responses, either using Salesforce's built-in models or custom ones.
- \* Test and validate the generated responses to ensure accuracy and effectiveness.
- \* Option B is correct as it follows the proper steps for using Prompt Builder.
- \* Option A and Option C do not capture the full process correctly.

References:

\* Salesforce Prompt Builder Documentation: [https://help.salesforce.com/s/articleView?id=sf.prompt\\_builder\\_overview.htm](https://help.salesforce.com/s/articleView?id=sf.prompt_builder_overview.htm)

**NO.33** Universal Containers (UC) wants to offer personalized service experiences and reduce agent handling time with AI-generated email responses, grounded in Knowledge base.

Which AI capability should UC use?

- A.** Einstein Email Replies
- B.** Einstein Service Replies for Email
- C.** Einstein Generative Service Replies for Email

**Answer:** B

Explanation:

For Universal Containers (UC) to offer personalized service experiences and reduce agent handling time using AI-generated responses grounded in the Knowledge base, the best solution is Einstein Service Replies for Email. This capability leverages AI to automatically generate responses to service-related emails based on historical data and the Knowledge base, ensuring accuracy and relevance while saving time for service agents.

- \* Einstein Email Replies (option A) is more suited for sales use cases.
- \* Einstein Generative Service Replies for Email (option C) could be a future offering, but as of now, Einstein Service Replies for Email is the correct choice for grounded, knowledge-based responses.

References:

\* Einstein Service Replies Overview: [https://help.salesforce.com/s/articleView?id=sf.einstein\\_service\\_replies.htm](https://help.salesforce.com/s/articleView?id=sf.einstein_service_replies.htm)

**NO.34** Universal Containers plans to implement prompt templates that utilize the standard foundation models.

What should the AI Specialist consider when building prompt templates in Prompt Builder?

- A.** Include multiple-choice questions within the prompt to test the LLM's understanding of the context.
- B.** Ask it to role-play as a character in the prompt template to provide more context to the LLM.

**C.** Train LLM with data using different writing styles including word choice, intensifiers, emojis, and punctuation.

**Answer:** B

Explanation:

When building prompt templates in Prompt Builder, it is essential to consider how the Large Language Model (LLM) processes and generates outputs. Training the LLM with various writing styles, such as different word choices, intensifiers, emojis, and punctuation, helps the model better understand diverse writing patterns and produce more contextually appropriate responses.

This approach enhances the flexibility and accuracy of the LLM when generating outputs for different use cases, as it is trained to recognize various writing conventions and styles. The prompt template should focus on providing rich context, and this stylistic variety helps improve the model's adaptability.

Options A and B are less relevant because adding multiple-choice questions or role-playing scenarios doesn't contribute significantly to improving the AI's output generation quality within standard business contexts.

For more details, refer to Salesforce's Prompt Builder documentation and LLM tuning strategies.

**NO.35** An AI specialist wants to leverage Record Snapshots grounding feature in a prompt template. What preparations are required?

- A.** Configure page layout of the master record type
- B.** Create a field set for all the fields to be grounded
- C.** Enable and configure dynamic form for the object

**Answer:** B

Explanation:

To use the Record Snapshots grounding feature in a prompt template, you must create a field set that includes all fields required for grounding. Field sets define which fields from an object are accessible to the AI model, ensuring the prompt template has structured data to generate contextually accurate responses.

Salesforce documentation emphasizes that grounding relies on explicitly defined field sets to avoid exposing unintended data and to comply with security policies.

\* Page Layout configuration (A) controls UI organization but does not directly enable grounding.

\* Dynamic Forms (C) customize record pages dynamically but are unrelated to data grounding for prompts.

**NO.36** Universal Containers (UC) wants to leverage Generative AI Salesforce functionality to reduce Service Agent handling time by providing recommended replies based on the existing Knowledge articles.

On which AI capability should UC train the service agents?

- A.** Case Replies
- B.** Knowledge Replies
- C.** Service Replies

**Answer:** B

Explanation:

Knowledge Replies is the Einstein Generative AI capability that generates recommended responses for service agents by grounding responses in the organization's Knowledge articles. This reduces

handling time by providing contextually accurate suggestions sourced directly from approved content. Salesforce documentation states that Knowledge Replies leverage natural language processing (NLP) to match customer inquiries with relevant articles and draft replies, ensuring consistency and compliance.

\* Case Replies (A) focus on generating responses based on case data (e.g., case fields, history) but do not explicitly ground responses in Knowledge articles.

\* Service Replies (C) is not a standard Einstein capability tied to Knowledge-driven responses.

**NO.37** Universal Containers needs to provide insights on the usability of Agents to drive adoption in the organization.

What should the AI Specialist recommend?

**A.** Agent Analytics

**B.** Agentforce Analytics

**C.** Agent Studio Analytics

**Answer:** A

Explanation:

\* Agent Analytics: This tool is specifically designed to provide usability insights for Salesforce agents. It tracks metrics like adoption rates, task completion times, and efficiency levels, helping organizations identify areas where agents excel or need additional support.

\* Agentforce Analytics: This term does not correspond to a recognized Salesforce feature.

\* Agent Studio Analytics: This is unrelated to analyzing agent usability, as it primarily supports customization or development features rather than providing analytics for adoption.

Thus, Agent Analytics is the correct recommendation as it offers actionable insights to drive agent adoption and productivity.

**NO.38** Which part of the Einstein Trust Layer architecture leverages an organization's own data within a large language model (LLM) prompt to confidently return relevant and accurate responses?

**A.** Prompt Defense

**B.** Data Masking

**C.** Dynamic Grounding

**Answer:** C

Explanation:

Dynamic Grounding in the Einstein Trust Layer architecture ensures that large language model (LLM) prompts are enriched with organization-specific data (e.g., Salesforce records, Knowledge articles) to generate accurate and relevant responses. By dynamically injecting contextual data into prompts, it reduces hallucinations and aligns outputs with trusted business data.

\* Prompt Defense (A) focuses on blocking malicious inputs or prompt injections but does not enhance responses with organizational data.

\* Data Masking (B) redacts sensitive information but does not contribute to grounding responses in business context.

**NO.39** Universal Containers (UC) is tracking web activities in Data Cloud for a unified contact, and wants to use that in a prompt template to help extract insights from the data.

Assuming that the Contact object is one of the objects associated with the prompt template, what is a valid way for DC to do this?

- A.** Call the prompt directly from Data Cloud with a web tracing activity included in the prompt definition.
- B.** Add the activity records as an enrichment related list to the Contact then pass the Contact into a prompt template workspace using related list grounding.
- C.** Create a prompt template that takes a list of all Data Cloud activity records as input to pass to the large language model (LLM).

**Answer:** B

Explanation:

To integrate web activity data from Data Cloud into a prompt template, the correct approach is to enrich the Contact object with the activity records as a related list and use related list grounding (Option B). Here's why:

- \* **Data Cloud Integration:** Data Cloud unifies web activity data and associates it with the unified Contact record. By adding these activities as a related list to the Contact, the data becomes accessible to the prompt template.
- \* **Prompt Template Grounding:** Salesforce prompt templates support grounding on related records. When the Contact is passed to the prompt template, the template can reference the related web activity records (via the related list) to extract insights.
- \* **Structured Data Handling:** This method aligns with Salesforce best practices for grounding, ensuring the large language model (LLM) receives structured, context-rich data without overwhelming it with raw activity lists.

Why Other Options Are Incorrect:

- \* **A. Calling the prompt directly from Data Cloud:** Prompt templates are invoked within Salesforce, not directly from Data Cloud. Grounding requires associating data with Salesforce objects, not ad-hoc web activity inclusion.
- \* **C. Passing a list of activity records as input:** While technically possible, this bypasses Salesforce's grounding framework, which relies on object relationships. It also risks exceeding LLM input limits and lacks scalability.

References:

- \* **Salesforce Data Cloud Implementation Guide:** Explains how to enrich standard/custom objects with related data for AI use cases.
- \* **Prompt Template Documentation:** Highlights grounding on related lists to leverage contextual data for LLM prompts.
- \* **Trailhead Module: "Einstein Prompt Builder Basics"** demonstrates grounding techniques using related records.

**NO.40** Universal Containers (UC) is using Einstein Generative AI to generate an account summary. UC aims to ensure the content is safe and inclusive, utilizing the Einstein Trust Layer's toxicity scoring to assess the content's safety level.

What does a safety category score of 1 indicate in the Einstein Generative Toxicity Score?

- A.** Not safe
- B.** Safe
- C.** Moderately safe

**Answer:** B

Explanation:

In the Einstein Trust Layer, the toxicity scoring system is used to evaluate the safety level of content

generated by AI, particularly to ensure that it is non-toxic, inclusive, and appropriate for business contexts. A toxicity score of 1 indicates that the content is deemed safe.

The scoring system ranges from 0 (unsafe) to 1 (safe), with intermediate values indicating varying degrees of safety. In this case, a score of 1 means that the generated content is fully safe and meets the trust and compliance guidelines set by the Einstein Trust Layer.

For further reference, check Salesforce's official Einstein Trust Layer documentation regarding toxicity scoring for AI-generated content.

**NO.41** Universal Containers is considering leveraging the Einstein Trust Layer in conjunction with Einstein Generative AI Audit Data.

Which audit data is available using the Einstein Trust Layer?

- A.** Response accuracy and offensiveness score
- B.** Hallucination score and bias score
- C.** Masked data and toxicity score

**Answer:** C

Explanation:

Universal Containers is considering the use of the Einstein Trust Layer along with Einstein Generative AI Audit Data. The Einstein Trust Layer provides a secure and compliant way to use AI by offering features like data masking and toxicity assessment.

The audit data available through the Einstein Trust Layer includes information about masked data-which ensures sensitive information is not exposed-and the toxicity score, which evaluates the generated content for inappropriate or harmful language.

References:

\* Salesforce AI Specialist Documentation - Einstein Trust Layer: Details the auditing capabilities, including logging of masked data and evaluation of generated responses for toxicity to maintain compliance and trust.

**NO.42** Amid their busy schedules, sales reps at Universal Containers dedicate time to follow up with prospects and existing clients via email regarding renewals or new deals. They spend many hours throughout the week reviewing past communications and details about their customers before performing their outreach.

Which standard Copilot action helps sales reps draft personalized emails to prospects by generating text based on previous successful communications?

- A.** Einstein Copilot Action: Find Similar Opportunities
- B.** Einstein Copilot Action: Draft or Revise Sales Email
- C.** Einstein Copilot Action: Summarize Record

**Answer:** B

Explanation:

For sales reps who need to draft personalized emails based on previous communications, the AI Specialist should recommend the Einstein Copilot Action: Draft or Revise Sales Email. This action uses AI to generate or revise email content, leveraging past successful communications to create personalized and relevant outreach to prospects or clients.

\* Find Similar Opportunities is used for opportunity matching, not email drafting.

\* Summarize Record provides a summary of customer data but does not directly help with drafting emails.



For more information, refer to Salesforce's Einstein Copilot documentation on standard actions for sales teams.

**NO.43** How does an Agent respond when it can't understand the request or find any requested information?

- A.** With a generated error message
- B.** With a general message asking the user to rephrase the request
- C.** With a preconfigured message, based on the action type

**Answer:** C

Explanation:

\* Context of the Question

\* When an Agent (e.g., Agentforce or a similar generative AI assistant in Salesforce) cannot understand a user's request or fails to locate the requested information, it should provide a coherent fallback.

\* Fallback / Error Handling in Agentforce

\* Preconfigured or "Fallback" Message: Typically, within the setup of any AI assistant, admins define a fallback or "failure" response. This message can be tailored to each action type or scenario (e.g., "No data found," "Sorry, I didn't get that-please try again," etc.).

\* Why Not an Automatically Generated Error Message? (Option A) AI assistants rarely show raw system error messages; they generally display friendly, standardized messages.

\* Why Not a Purely Generic "Rephrase" Message? (Option B) Although an agent might prompt the user to rephrase, Salesforce best practices are to configure a fallback response or fallback action that aligns with the brand and the context. This is typically a "preconfigured message based on the action type."

\* Conclusion Because the assistant uses a preconfigured fallback that is set up in the environment (for example, in the agent's or domain's settings), the correct choice is Option C.

Salesforce AI Specialist References & Documents

\* Salesforce Pilot / Agentforce Setup Documentation Explains how to configure a fallback or default message when the AI cannot fulfill a user's request.

\* Salesforce AI Specialist Study Guide Details best practices for AI-driven assistants and how fallback scenarios are handled with preconfigured messages.

**NO.44** Universal Containers (UC) wants to improve the efficiency of addressing customer questions and reduce agent handling time with AI-generated responses. The agents should be able to leverage their existing knowledge base and identify whether the responses are coming from the large language model (LLM) or from Salesforce Knowledge.

Which step should UC take to meet this requirement?

- A.** Turn on Service AI Grounding, Grounding with Case, and Service Replies.
- B.** Turn on Service Replies, Service AI Grounding, and Grounding with Knowledge.
- C.** Turn on Service AI Grounding and Grounding with Knowledge.

**Answer:** C

Explanation:

To meet Universal Containers' goal of improving efficiency and reducing agent handling time with AI-generated responses, the best approach is to enable Service Replies, Service AI Grounding, and Grounding with Knowledge.

- \* Service Replies generates responses automatically.
  - \* Service AI Grounding ensures that the AI is using relevant case data.
  - \* Grounding with Knowledge ensures that responses are backed by Salesforce Knowledge articles, allowing agents to identify whether a response is coming from the LLM or Salesforce Knowledge.
  - \* Option C does not include Service Replies, which is necessary for generating AI responses.
  - \* Option A lacks the Grounding with Knowledge, which is essential for identifying response sources.
- For more details, refer to Salesforce Service AI documentation on grounding and service replies.

**NO.45** Universal Containers (UC) noticed an increase in customer contract cancellations in the last few months. UC is seeking ways to address this issue by implementing a proactive outreach program to customers before they cancel their contracts and is asking the Salesforce team to provide suggestions.

Which use case functionality of Model Builder aligns with UC's request?

- A.** Product recommendation prediction
- B.** Customer churn prediction
- C.** Contract Renewal Date prediction

**Answer:** B

Explanation:

Customer churn prediction is the best use case for Model Builder in addressing Universal Containers' concerns about increasing customer contract cancellations. By implementing a model that predicts customer churn, UC can proactively identify customers who are at risk of canceling and take action to retain them before they decide to terminate their contracts. This functionality allows the business to forecast churn probability based on historical data and initiate timely outreach programs.

- \* Option B is correct because customer churn prediction aligns with UC's need to reduce cancellations through proactive measures.
- \* Option A (product recommendation prediction) is unrelated to contract cancellations.
- \* Option C (contract renewal date prediction) addresses timing but does not focus on predicting potential cancellations.

References:

- \* Salesforce Model Builder Use Case Overview: [https://help.salesforce.com/s/articleView?id=sf.model\\_builder\\_use\\_cases.htm](https://help.salesforce.com/s/articleView?id=sf.model_builder_use_cases.htm)

**NO.46** After a successful implementation of Agentforce Sales Agent with sales users. Universal Containers now aims to deploy it to the service team.

Which key consideration should the AI Specialist keep in mind for this deployment?

- A.** Assign the Agentforce for Service permission to the Service Cloud users.
- B.** Assign the standard service actions to Agentforce Service Agent.
- C.** Review and test standard and custom Agent topics and actions for Service Center use cases.

**Answer:** C

Explanation:

When deploying Einstein Agent (formerly Agentforce) from Sales to Service Cloud:

- \* Agent Topics and Actions are context-specific. Service Cloud use cases (e.g., case resolution, knowledge retrieval) require validation of existing topics/actions to ensure alignment with service workflows.
- \* Option A: Permissions like "Agentforce for Service" are necessary but secondary to functional

compatibility.

\* Option B: Standard service actions must be mapped to Agentforce, but testing ensures they function as intended.

References:

\* Salesforce Help: Einstein Agent Setup

\* Emphasizes reviewing "topics and actions for different user groups (Sales vs. Service)."

**NO.47** What is the role of the large language model (LLM) in understanding intent and executing an Agent Action?

**A.** Identify the best matching topic and actions and correct order of execution

**B.** Determine a user's topic access and sort actions by priority to be executed

**C.** Find similar requested topics and provide the actions that need to be executed

**Answer: A**

**NO.48** A data science team has trained an XGBoost classification model for product recommendations on Databricks. The AI Specialist is tasked with bringing inferences for product recommendations from this model into Data Cloud as a stand-alone data model object (DMO). How should the AI Specialist set this up?

**A.** Create the serving endpoint in Databricks, then configure the model using Model Builder.

**B.** Create the serving endpoint in Einstein Studio, then configure the model using Model Builder.

**C.** Create the serving endpoint in Databricks, then configure the model using a Python SDK connector.

**Answer: A**

Explanation:

To integrate inferences from an XGBoost model into Salesforce's Data Cloud as a stand-alone Data Model Object (DMO):

\* Create the Serving Endpoint in Databricks:

\* The serving endpoint is necessary to make the trained model available for real-time inference. Databricks provides tools to host and expose the model via an endpoint.

\* Configure the Model Using Model Builder:

\* After creating the endpoint, the AI Specialist should configure it within Einstein Studio's Model Builder, which integrates external endpoints with Salesforce Data Cloud for processing and storing inferences as DMOs.

\* Option B: Serving endpoints are not created in Einstein Studio; they are set up in external platforms like Databricks before integration.

\* Option C: A Python SDK connector is not used to bring model inferences into Salesforce Data Cloud; Model Builder is the correct tool.

**NO.49** A service agent is looking at a custom object that stores travel information. They recently received a weather alert and now need to cancel flights for the customers that are related with this itinerary. The service agent needs to review the Knowledge articles about canceling and rebooking the customer flights.

Which Einstein Copilot capability helps the agent accomplish this?

**A.** Execute tasks based on available actions, answering questions using information from accessible Knowledge articles.

- B.** Invoke a flow which makes a call to external data to create a Knowledge article.
- C.** Generate a Knowledge article based off the prompts that the agent enters to create steps to cancel flights.

**Answer:** C

Explanation:

In this scenario, the Einstein Copilot capability that best helps the agent is its ability to execute tasks based on available actions and answer questions using data from Knowledge articles. Einstein Copilot can assist the service agent by providing relevant Knowledge articles on canceling and rebooking flights, ensuring that the agent has access to the correct steps and procedures directly within the workflow.

This feature leverages the agent's existing context (the travel itinerary) and provides actionable insights or next steps from the relevant Knowledge articles to help the agent quickly resolve the customer's needs.

The other options are incorrect:

- \* B refers to invoking a flow to create a Knowledge article, which is unrelated to the task of retrieving existing Knowledge articles.
- \* C focuses on generating Knowledge articles, which is not the immediate need for this situation where the agent requires guidance on existing procedures.

References:

- \* Salesforce Documentation on Einstein Copilot
- \* Trailhead Module on Einstein for Service

**NO.50** In Model Playground, which hyperparameters of an existing Salesforce-enabled foundational model can an AI Specialist change?

- A.** Temperature, Frequency Penalty, Presence Penalty
- B.** Temperature, Top-k sampling, Presence Penalty
- C.** Temperature, Frequency Penalty, Output Tokens

**Answer:** A

Explanation:

In Model Playground, an AI specialist working with a Salesforce-enabled foundational model has control over specific hyperparameters that can directly affect the behavior of the generative model:

- \* Temperature: Controls the randomness of predictions. A higher temperature leads to more diverse outputs, while a lower temperature makes the model's responses more focused and deterministic.
- \* Frequency Penalty: Reduces the likelihood of the model repeating the same phrases or outputs frequently.
- \* Presence Penalty: Encourages the model to introduce new topics in its responses, rather than sticking with familiar, previously mentioned content.

These hyperparameters are adjustable to fine-tune the model's responses, ensuring that it meets the desired behavior and use case requirements. Salesforce documentation confirms that these three are the key tunable hyperparameters in the Model Playground.

For more details, refer to Salesforce AI Model Playground guidance from Salesforce's official documentation on foundational model adjustments.

**NO.51** Universal Containers wants to incorporate the current order fulfillment status into a prompt for a large language model (LLM). The order status is stored in the external enterprise resource

planning (ERP) system.

Which data grounding technique should the AI Specialist recommend?

- A.** Eternal Object Record Merge Fields
- B.** External Services Merge Fields
- C.** Apex Merge Fields

**Answer:** A

- \* Context of the Requirement: Universal Containers wants to pull in real-time order status data from an external ERP system into an LLM prompt.
- \* Data Grounding in LLM Prompts: Data grounding ensures the Large Language Model has access to the most current and relevant information. In Salesforce, one recommended approach is to use External Objects (via Salesforce Connect) when data resides outside of Salesforce.
- \* Why External Object Record Merge Fields:
  - \* External Objects appear much like standard or custom objects but map to tables in external systems.
  - \* You can reference fields from these External Objects in merge fields, allowing real-time data retrieval from the external ERP system without storing that data natively in Salesforce.
  - \* This is a simpler "point-and-reference" approach compared to coding custom Apex or configuring external services for direct prompt embedding.
- \* Why Not External Services Merge Fields or Apex Merge Fields:
  - \* External Services Merge Fields typically leverage flows or external service definitions. While feasible, it is more about orchestrating or invoking external services for automation (e.g., Flow). It's not the standard approach for seamlessly referencing external record data in prompt merges.
  - \* Apex Merge Fields would imply custom Apex code controlling the prompt insertion. While possible, it's less "clicks not code" friendly and is not the default method for referencing typical record data.
- \* References and Study Resources:
  - \* Salesforce Help & Training # Salesforce Connect and External Objects
  - \* Salesforce Trailhead # "Integrate External Data with Salesforce Connect"
  - \* Salesforce AI Specialist Study Resources (documentation regarding how to ground LLM prompts using External Objects)

**NO.52** Universal Containers aims to streamline the sales team's daily tasks by using AI.

When considering these new workflows, which improvement requires the use of Prompt Builder?

- A.** Populate an AI-generated time-to close estimation to opportunities
- B.** Populate an AI generated summary field for sales contracts.
- C.** Populate an AI generated lead score for new leads.

**Answer:** B

Explanation:

Prompt Builder is explicitly required to create AI-generated summary fields via prompt templates. These fields use natural language instructions to extract or synthesize information (e.g., summarizing contract terms). Time-to-close estimations (A) and lead scores (C) are typically handled by predictive AI (e.g., Einstein Opportunity Scoring) or analytics tools, which do not require Prompt Builder.

**NO.53** An AI Specialist created a custom Agent action, but it is not being picked up by the planner service in the correct order.

Which adjustment should the AI Specialist make in the custom Agent action instructions for the

planner service to work as expected?

- A.** Specify the dependent actions with the reference to the action API name.
- B.** Specify the profiles or custom permissions allowed to invoke the action.
- C.** Specify the LLM model provider and version to be used to invoke the action.

**Answer:** A

Explanation:

When a custom Agent action is not being prioritized correctly by the planner service, the root cause is often missing or improperly defined action dependencies. The planner service determines the execution order of actions based on dependencies defined in the action instructions. To resolve this, the AI Specialist must explicitly specify dependent actions using their API names in the custom action's configuration. This ensures the planner understands the sequence in which actions must be executed to meet business logic requirements.

Salesforce documentation highlights that dependencies are critical for orchestrating workflows in Einstein Bots and Agentforce. For example, if Action B requires data from Action A, Action A's API name must be listed as a dependency in Action B's instructions. The Einstein Bot Developer Guide states that failing to define dependencies can lead to race conditions or incorrect execution order.

In contrast:

- \* Profiles or custom permissions (B) control access to the action but do not influence execution order.
- \* LLM model provider and version (C) determine the AI model used for processing but are unrelated to the planner's sequencing logic.

**NO.54** Universal Containers (UC) plans to automatically populate the Description field on the Account object.

Which type of prompt template should UC use?

- A.** Field Generation prompt template
- B.** Flex Prompt template
- C.** Sales Email prompt template

**Answer:** A

Explanation:

\* Context of the Question Universal Containers (UC) wants to automatically populate the Description field on the Account object. The AI-driven solution must generate textual data and write it directly into a field.

\* Field Generation Prompt Template

\* Primary Use Case: A Field Generation prompt template is specifically designed to create or fill in fields on a record with AI-generated text.

\* Auto-population: By configuring a Field Generation prompt template, admins can define the instructions, data inputs, and desired output for the AI. The resulting text then populates the specified field, such as the Account Description.

\* Why Not Flex or Sales Email Prompt Templates?

\* Flex Prompt Template: Used to combine or manipulate data across objects, merges, or references from multiple sources in more advanced, flexible prompts. Typically not the go-to for straightforward text generation on a single field.

\* Sales Email Prompt Template: Focused on drafting or summarizing emails for sales reps (like crafting outreach or follow-up messages). This template is not specifically built to populate a field on

a record.

\* Conclusion For automatically populating the Description field with AI-generated content, the Field Generation prompt template (Option A) is the correct choice.

Salesforce AI Specialist References & Documents

\* Salesforce Documentation: Prompt Template Types Explains various template types (Field Generation, Flex, Email, etc.) and their typical use cases.

\* Salesforce AI Specialist Study Guide Highlights Field Generation prompt templates for populating or updating record fields with AI-generated text.

**NO.55** What is the main purpose of Prompt Builder?

**A.** A tool for developers to use in Visual Studio Code that creates prompts for Apex programming, assisting developers in writing code more efficiently.

**B.** A tool that enables companies to create reusable prompts for large language models (LLMs), bringing generative AI responses to their flow of work

**C.** A tool within Salesforce offering real-time AI-powered suggestions and guidance to users, Improving productivity and decision-making.

**Answer:** B

Explanation:

Prompt Builder is designed to help organizations create and configure reusable prompts for large language models (LLMs). By integrating generative AI responses into workflows, Prompt Builder enables customization of AI prompts that interact with Salesforce data and automate complex processes. This tool is especially useful for creating tailored and consistent AI-generated content in various business contexts, including customer service and sales.

\* It is not a tool for Apex programming (as in option A).

\* It is also not limited to real-time suggestions as mentioned in option C. Instead, it provides a flexible way for companies to manage and customize how AI-driven responses are generated and used in their workflows.

References:

\* Salesforce Prompt Builder Overview:

[https://help.salesforce.com/s/articleView?id=sf.prompt\\_builder.htm](https://help.salesforce.com/s/articleView?id=sf.prompt_builder.htm)

**NO.56** Universal Containers plans to enhance the customer support team's productivity using AI. Which specific use case necessitates the use of Prompt Builder?

**A.** Creating a draft of a support bulletin post for new product patches

**B.** Creating an AI-generated customer support agent performance score

**C.** Estimating support ticket volume based on historical data and seasonal trends

**Answer:** A

Explanation:

The use case that necessitates the use of Prompt Builder is creating a draft of a support bulletin post for new product patches. Prompt Builder allows the AI Specialist to create and refine prompts that generate specific, relevant outputs, such as drafting support communication based on product information and patch details.

\* Option B (agent performance score) would likely involve predictive modeling, not prompt generation.

\* Option C (estimating support ticket volume) would require data analysis and predictive tools, not prompt building.

For more details, refer to Salesforce's Prompt Builder documentation for generative AI content creation.

**NO.57** Universal Containers has seen a high adoption rate of a new feature that uses generative AI to populate a summary field of a custom object, Competitor Analysis. All sales users have the same profile but one user cannot see the generative AI-enabled field icon next to the summary field. What is the most likely cause of the issue?

- A.** The user does not have the Prompt Template User permission set assigned.
- B.** The prompt template associated with summary field is not activated for that user.
- C.** The user does not have the field Generative AI User permission set assigned.

**Answer: C**

Explanation:

In Salesforce, Generative AI capabilities are controlled by specific permission sets. To use features such as generating summaries with AI, users need to have the correct permission sets that allow access to these functionalities.

\* Generative AI User Permission Set: This is a key permission set required to enable the generative AI capabilities for a user. In this case, the missing Generative AI User permission set prevents the user from seeing the generative AI-enabled field icon. Without this permission, the generative AI feature in the Competitor Analysis custom object won't be accessible.

\* Why not A? The Prompt Template User permission set relates specifically to users who need access to prompt templates for interacting with Einstein GPT, but it's not directly related to the visibility of AI-enabled field icons.

\* Why not B? While a prompt template might need to be activated, this is not the primary issue here. The question states that other users with the same profile can see the icon, so the problem is more likely to be permissions-based for this particular user.

For more detailed information, you can review Salesforce documentation on permission sets related to AI capabilities at [Salesforce AI Documentation](#) and [Einstein GPT permissioning guidelines](#).

**NO.58** Universal Containers wants to implement a solution in Salesforce with a custom UX that allows users to enter a sales order number.

Subsequently, the system will invoke a custom prompt template to create and display a summary of the sales order header and sales order details.

Which solution should an AI Specialist implement to meet this requirement?

- A.** Create a screen flow to collect sales order number and invoke the prompt template using the standard "Prompt Template" flow action.
- B.** Create a template-triggered prompt flow and invoke the prompt template using the standard "Prompt Template" flow action.
- C.** Create an autolaunched flow and invoke the prompt template using the standard "Prompt Template" flow action.

**Answer: A**

Explanation:

To implement a solution where users enter a sales order number and the system generates a



summary, the AI Specialist should create a screen flow to collect the sales order number and invoke the prompt template. The standard "Prompt Template" flow action can then be used to trigger the custom prompt, providing a summary of the sales order header and details.

\* Option B, creating a template-triggered prompt flow, is not necessary for this scenario because the requirement is to directly collect input through a screen flow.

\* Option C, using an autolaunched flow, would be inappropriate here because the solution requires user interaction (entering a sales order number), which is best suited to a screen flow.

Salesforce AI Specialist References: For further guidance on creating prompt templates with flows: [https://help.salesforce.com/s/articleView?id=sf.prompt\\_template\\_flow\\_integration.htm](https://help.salesforce.com/s/articleView?id=sf.prompt_template_flow_integration.htm)

**NO.59** How should an organization use the Einstein Trust layer to audit, track, and view masked data?

**A.** Utilize the audit trail that captures and stores all LLM submitted prompts in Data Cloud.

**B.** In Setup, use Prompt Builder to send a prompt to the LLM requesting for the masked data.

**C.** Access the audit trail in Setup and export all user-generated prompts.

**Answer:** A

Explanation:

The Einstein Trust Layer is designed to ensure transparency, compliance, and security for organizations leveraging Salesforce's AI and generative AI capabilities. Specifically, for auditing, tracking, and viewing masked data, organizations can utilize:

\* Audit Trail in Data Cloud: The audit trail captures and stores all prompts submitted to large language models (LLMs), ensuring that sensitive or masked data interactions are logged. This allows organizations to monitor and audit all AI-generated outputs, ensuring that data handling complies with internal and regulatory guidelines. The Data Cloud provides the infrastructure for managing and accessing this audit data.

\* Why not B? Using Prompt Builder in Setup to send prompts to the LLM is for creating and managing prompts, not for auditing or tracking data. It does not interact directly with the audit trail functionality.

\* Why not C? Although the audit trail can be accessed in Setup, the user-generated prompts are primarily tracked in the Data Cloud for broader control, auditing, and analysis. Setup is not the primary tool for exporting or managing these audit logs.

More information on auditing AI interactions can be found in the Salesforce AI Trust Layer documentation, which outlines how organizations can manage and track generative AI interactions securely.

**NO.60** Universal Containers' current AI data masking rules do not align with organizational privacy and security policies and requirements.

What should an AI Specialist recommend to resolve the issue?

**A.** Enable data masking for sandbox refreshes.

**B.** Configure data masking in the Einstein Trust Layer setup.

**C.** Add new data masking rules in LLM setup.

**Answer:** B

Explanation:

When Universal Containers' AI data masking rules do not meet organizational privacy and security standards, the AI Specialist should configure the data masking rules within the Einstein Trust Layer.

The Einstein Trust Layer provides a secure and compliant environment where sensitive data can be masked or anonymized to adhere to privacy policies and regulations.

\* Option A, enabling data masking for sandbox refreshes, is related to sandbox environments, which are separate from how AI interacts with production data.

\* Option C, adding masking rules in the LLM setup, is not appropriate because data masking is managed through the Einstein Trust Layer, not the LLM configuration.

The Einstein Trust Layer allows for more granular control over what data is exposed to the AI model and ensures compliance with privacy regulations.

Salesforce AI Specialist References: For more information, refer to: [https://help.salesforce.com/s/articleView?id=sf.einstein\\_trust\\_layer\\_data\\_masking.htm](https://help.salesforce.com/s/articleView?id=sf.einstein_trust_layer_data_masking.htm)

**NO.61** Universal Containers (UC) is looking to enhance its operational efficiency. UC has recently adopted Salesforce and is considering implementing Einstein Copilot to improve its processes.

What is a key reason for implementing Einstein Copilot?

- A.** Improving data entry and data cleansing
- B.** Allowing AI to perform tasks without user interaction
- C.** Streamlining workflows and automating repetitive tasks

**Answer:** C

Explanation:

The key reason for implementing Einstein Copilot is its ability to streamline workflows and automate repetitive tasks. By leveraging AI, Einstein Copilot can assist users in handling mundane, repetitive processes, such as automatically generating insights, completing actions, and guiding users through complex processes, all of which significantly improve operational efficiency.

\* Option A (Improving data entry and cleansing) is not the primary purpose of Einstein Copilot, as its focus is on guiding and assisting users through workflows.

\* Option B (Allowing AI to perform tasks without user interaction) does not accurately describe the role of Einstein Copilot, which operates interactively to assist users in real time.

Salesforce AI Specialist References: More details can be found in the Salesforce documentation: [https://help.salesforce.com/s/articleView?id=sf.einstein\\_copilot\\_overview.htm](https://help.salesforce.com/s/articleView?id=sf.einstein_copilot_overview.htm)

**NO.62** Universal Container (UC) has effectively utilized prompt templates to update summary fields on Lightning record pages. An admin now wishes to incorporate similar functionality into UC's automation process using Flow.

How can the admin get a response from this prompt template from within a flow to use as part of UC's automation?

- A.** Invocable Apex
- B.** Flow Action
- C.** Einstein for Flow

**Answer:** C

Explanation:

\* Context of the Question

\* Universal Container (UC) has used prompt templates to update summary fields on record pages.

\* Now, the admin wants to incorporate similar generative AI functionality within a Flow for automation purposes.

\* How to Call a Prompt Template Within a Flow

\* Flow Action: Salesforce provides a standard way to invoke generative AI templates or prompts within a Flow step. From the Flow Builder, you can add an "Action" that references the prompt template you created in Prompt Builder.

\* Other Options:

\* Invocable Apex: Possible fallback if there's no out-of-the-box Flow Action available.

However, Salesforce is releasing native Flow integration for AI prompts, making custom Apex less necessary.

\* Einstein for Flow: A broad label for Salesforce's generative AI features within Flow.

Under the hood, you typically use a "Flow Action" that points to your prompt.

\* Conclusion

\* The easiest out-of-the-box solution is to use a Flow Action referencing the prompt template.

Hence, Option B is correct.

Salesforce AI Specialist References & Documents

\* Salesforce Trailhead: Use Prompt Templates in Flow Demonstrates how to add an Action in Flow that calls a prompt template.

\* Salesforce Documentation: Einstein GPT for Flow Explains standard flow actions to invoke and handle generative AI responses.

**NO.63** What does it mean when a prompt template version is described as immutable?

**A.** Only the latest version of a template can be activated.

**B.** Every modification on a template will be saved as a new version automatically.

**C.** Prompt template version is activated; no further changes can be saved to that version.

**Answer:** C

Explanation:

When a prompt template version is immutable, it means that once the version is activated, it cannot be edited or modified. This ensures consistency in production environments where changes could disrupt workflows.

\* Option A is incorrect: Any version (not just the latest) can be activated, depending on the use case.

\* Option D is incorrect: Modifications require manually creating a new version; automatic versioning is not enforced.

\* Option C is correct: Activation locks the version, enforcing immutability.

References:

\* Salesforce Help: Prompt Template Versioning

\* States that "activated prompt template versions are immutable and cannot be edited."

**NO.64** When configuring a prompt template, an AI Specialist previews the results of the prompt template they've written. They see two distinct text outputs: Resolution and Response. Which information does the Resolution text provide?

**A.** It shows the full text that is sent to the Trust Layer.

**B.** It shows the response from the LLM based on the sample record.

**C.** It shows which sensitive data is masked before it is sent to the LLM.

**Answer:** B

Explanation:

When previewing a prompt template in Salesforce, the Resolution text provides the response from

the LLM (Large Language Model) based on the data from a sample record. This output shows what the AI model generated in response to the prompt, giving the AI Specialist a chance to review and adjust the response before finalizing the template.

\* Option B is correct because Resolution displays the actual response generated by the LLM.

\* Option A refers to sending the text to the Trust Layer, but that's not what Resolution represents.

\* Option C relates to data masking, which is shown elsewhere, not under Resolution.

References:

\* Salesforce Prompt Builder Overview: [https://help.salesforce.com/s/articleView?id=sf.prompt\\_builder\\_overview.htm](https://help.salesforce.com/s/articleView?id=sf.prompt_builder_overview.htm)

**NO.65** An AI Specialist needs to create a Sales Email with a custom prompt template. They need to ground on the following data.

Opportunity Products Events near the customer Tone and voice examples

How should the AI Specialist obtain related items?

**A.** Call prompt initiated flow to fetch and ground the required data.

**B.** Create a flex template that takes the records in question as inputs.

**C.** Utilize a standard email template and manually insert the required data fields.

**Answer:** A

Explanation:

To ground a sales email on Opportunity Products, Events near the customer, and Tone and voice examples, the AI Specialist should use a prompt-initiated flow. This flow can dynamically fetch the necessary data from related records in Salesforce and ground the generative AI output with contextually accurate information.

\* Option B (flex template) does not provide the ability to fetch dynamic data from Salesforce records automatically.

\* Option C (manual insertion) would not allow for the dynamic and automated grounding of data required for custom prompts.

Refer to Salesforce documentation on flows and grounding for more details on integrating data into custom prompt templates.

**NO.66** Universal Containers, dealing with a high volume of chat inquiries, implements Einstein Work Summaries to boost productivity.

After an agent-customer conversation, which additional information does Einstein generate and fill, apart from the "summary"?

**A.** Sentiment Analysis and Emotion Detection

**B.** Draft Survey Request Email

**C.** Issue and Resolution

**Answer:** C

Explanation:

Einstein Work Summaries automatically generate concise summaries of customer interactions (e.g., chat transcripts). Beyond the "summary" field, it extracts and populates Issue (key problem discussed) and Resolution (action taken to resolve the issue). These fields help agents and supervisors quickly grasp the conversation's context without reviewing the full transcript.

\* Sentiment Analysis and Emotion Detection (Option A): While Einstein Conversation Insights provides sentiment scores and emotion detection, these are separate from Work Summaries. Work

Summaries focus on factual summaries, not sentiment.

- \* Draft Survey Request Email (Option B): Not part of Work Summaries. This would require automation tools like Flow or Email Studio.
- \* Issue and Resolution (Option C): Directly referenced in Salesforce documentation as fields populated by Einstein Work Summaries.

References:

- \* Salesforce Help Article: Einstein Work Summaries
- \* Einstein Work Summaries focus on "key details like Issue and Resolution" alongside summaries.
- \* Contrast with Einstein Conversation Insights for sentiment/emotion analysis.

**NO.67** Universal Containers (UC) is discussing its AI strategy in an agile Scrum meeting.

Which business requirement would lead an AI Specialist to recommend connecting to an external foundational model via Einstein Studio (Model Builder)?

- A.** UC wants to fine-tune model temperature.
- B.** UC wants a model fine-tuned using company data.
- C.** UC wants to change the frequency penalty of the model.

**Answer: B**

Explanation:

Einstein Studio (Model Builder) allows organizations to connect and utilize external foundational models while fine-tuning them with company-specific data. This capability is particularly suited to businesses like Universal Containers (UC) that require customization of foundational models to better align with their unique data and use cases.

- \* Option A: Adjusting model temperature is a parameter-level setting for controlling randomness in AI-generated responses but does not necessitate connecting to an external foundational model.
- \* Option B: This is the correct answer because Einstein Studio supports fine-tuning external models with proprietary company data, enabling a tailored and more accurate AI solution for UC.
- \* Option C: Changing frequency penalties is another parameter-level adjustment and does not require external foundational models or Einstein Studio.

**NO.68** Universal Containers is very concerned about security compliance and wants to understand: Which prompt text is sent to the large language model (LLM)

- \* How it is masked
- \* The masked response

What should the AI Specialist recommend?

- A.** Ingest the Einstein Shield Event logs into CRM Analytics.
- B.** Review the debug logs of the running user.
- C.** Enable audit trail in the Einstein Trust Layer.

**Answer: C**

Explanation:

To address security compliance concerns and provide visibility into the prompt text sent to the LLM, how it is masked, and the masked response, the AI Specialist should recommend enabling the audit trail in the Einstein Trust Layer. This feature captures and logs the prompts sent to the large language model (LLM) along with the masking of sensitive information and the AI's response. This audit trail ensures full transparency and compliance with security requirements.

- \* Option A (Einstein Shield Event logs) is focused on system events rather than specific AI prompt

data.

\* Option B (debug logs) would not provide the necessary insight into AI prompt masking or responses.

For further details, refer to Salesforce's Einstein Trust Layer documentation about auditing and security measures.

**NO.69** An AI Specialist is considering using a Field Generation prompt template type.

What should the AI Specialist check before creating the Field Generation prompt to ensure it is possible for the field to be enabled for generative AI?

**A.** That the field chosen must be a rich text field with 255 characters or more.

**B.** That the org is set to API version 59 or higher

**C.** That the Lightning page layout where the field will reside has been upgraded to Dynamic Forms

**Answer: B**

Explanation:

Before creating a Field Generation prompt template, the AI Specialist must ensure that the Salesforce org is set to API version 59 or higher. This version of the API introduces support for advanced generative AI features, such as enabling fields for generative AI outputs. This is a critical technical requirement for the Field Generation prompt template to function correctly.

\* Option A (rich text field requirement) is not necessary for generative AI functionality.

\* Option C (Dynamic Forms) does not impact the ability of a field to be generative AI-enabled, although it might enhance the user interface.

For more information, refer to Salesforce documentation on API versioning and Field Generation templates.

**NO.70** Universal Containers (UC) recently rolled out Einstein Generative capabilities and has created a custom prompt to summarize case records. Users have reported that the case summaries generated are not returning the appropriate information.

What is a possible explanation for the poor prompt performance?

**A.** The data being used for grounding is incorrect or incomplete.

**B.** The prompt template version is incompatible with the chosen LLM.

**C.** The Einstein Trust Layer is incorrectly configured.

**Answer: A**

Explanation:

Poor prompt performance when generating case summaries is often due to the data used for grounding being incorrect or incomplete. Grounding involves feeding accurate, relevant data to the AI so it can generate appropriate outputs. If the data source is incomplete or contains errors, the generated summaries will reflect that by being inaccurate or insufficient.

\* Option B (prompt template incompatibility with the LLM) is unlikely because such incompatibility usually results in more technical failures, not poor content quality.

\* Option C (Einstein Trust Layer misconfiguration) is focused on data security and auditing, not the quality of prompt responses.

For more information, refer to Salesforce documentation on grounding AI models and data quality best practices.

**NO.71** Universal Containers' service team wants to customize the standard case summary response

from Einstein Copilot.

What should the AI Specialist do to achieve this?

- A.** Customize the standard Record Summary template for the Case object,
- B.** Summarize the Case with a standard copilot action.
- C.** Create a custom Record Summary prompt template for the Case object.

**Answer:** C

Explanation:

To customize the case summary response from Einstein Copilot, the AI Specialist should create a custom Record Summary prompt template for the Case object. This allows Universal Containers to tailor the way case data is summarized, ensuring the output aligns with specific business requirements or user preferences.

\* Option A (customizing the standard Record Summary template) does not provide the flexibility required for deep customization.

\* Option B (standard Copilot action) won't allow customization; it will only use default settings.

Refer to Salesforce Prompt Builder documentation for guidance on creating custom templates for record summaries.

**NO.72** Universal Containers is interested in using Call Explorer to quickly gain insights from meetings recorded by its sales team.

What should the AI Specialist be aware of before enabling this feature?

- A.** Call Explorer operates independently of Salesforce Knowledge, requiring no prior setup.
- B.** Custom Call Explorer actions need to be built before it can be configured.
- C.** Call Explorer requires the Einstein Conversation Insights permission set to be enabled.

**Answer:** C

Explanation:

Before enabling Call Explorer, the Salesforce AI Specialist must ensure that the Einstein Conversation Insights permission set is assigned to users (Option C). Call Explorer is a feature within Einstein Conversation Insights (ECI) that analyzes meeting recordings to surface trends, keywords, and actionable insights.

Key Considerations:

\* Permission Set Requirement:

\* Users (including admins) need the Einstein Conversation Insights permission set to access and use Call Explorer. Without this, the feature remains inaccessible.

\* The permission set grants access to ECI tools, including call transcription, analysis, and dashboard visibility.

\* Why Other Options Are Incorrect:

\* A. Independence from Salesforce Knowledge: While Call Explorer does not rely on Salesforce Knowledge, this is irrelevant to the setup prerequisite. The critical dependency is the permission set, not Knowledge configuration.

\* B. Custom Actions: Call Explorer does not require custom actions to be built before configuration. It is a pre-built analytics tool that works once permissions and data sources (e.g., call recordings) are configured.

References:

\* Salesforce Einstein Conversation Insights Guide: Explicitly states that the Einstein Conversation Insights permission set is required to access Call Explorer.

\* Trailhead Module: "Einstein Conversation Insights Basics" outlines permission prerequisites for enabling call analytics.

\* Salesforce Help Documentation: Confirms that Call Explorer functionality is governed by ECI permissions.

**NO.73** Universal Containers (UC) is using Einstein Generative AI to generate an account summary. UC aims to ensure the content is safe and inclusive, utilizing the Einstein Trust Layer's toxicity scoring to assess the content's safety level.

In the score of 1 indicate?

**A.** The response is the least toxic Einstein Generative AI Toxicity Scoring system, what does a toxicity category.

**B.** The response is not toxic.

**C.** The response is the most toxic.

**Answer:** C

Explanation:

Einstein Trust Layer's Toxicity Scoring categorizes content on a scale of 0 to 1, where 1 indicates the highest level of toxicity (e.g., harmful, biased, or inappropriate language). This scoring helps organizations filter unsafe AI-generated content. A score of 1 triggers mitigation actions, such as blocking the response or alerting administrators.

\* A score of 0 would indicate no toxicity (B is incorrect).

\* The scoring system does not use "least toxic" as a category (A is misleading).

**NO.74** An AI Specialist is tasked with creating a prompt template for a sales team. The template needs to generate a summary of all related opportunities for a given Account.

Which grounding technique should the AI Specialist use to include data from the related list of opportunities in the prompt template?

**A.** Use the merge fields to reference a custom related list of opportunities.

**B.** Use merge fields to reference the default related list of opportunities.

**C.** Use formula fields to reference the Einstein related list of opportunities.

**Answer:** B

Explanation:

In Salesforce, when creating a prompt template for the sales team, you can include data from related objects such as Opportunities that are linked to an Account. The best method to ground the AI model and provide relevant information from related records, like Opportunities, is by using merge fields. Merge fields in Salesforce allow you to dynamically reference data from a record or related records, like Opportunities for a given Account. In this scenario, the AI Specialist needs to pull data from the default related list of Opportunities associated with the Account. This is achieved by using merge fields, which pull in data from the standard relationship Salesforce creates between Accounts and Opportunities.

Option A (referencing a custom related list) and Option C (using formula fields with Einstein-related lists) do not align with the standard, practical grounding method for this task. Custom lists would require additional configurations not typically necessary for a basic use case, and formula fields are typically not used to directly fetch related list data for prompt generation in templates. The standard and straightforward method is using merge fields tied to the default related list of opportunities.

Salesforce References:



\* Merge Fields in Templates: <https://help.salesforce.com/s/articleView?id=000387601&type=1>

\* Grounding Data in Prompts: [https://developer.salesforce.com/docs/atlas.en-us.salesforce\\_ai.meta/salesforce\\_ai/grounding\\_data\\_prompts](https://developer.salesforce.com/docs/atlas.en-us.salesforce_ai.meta/salesforce_ai/grounding_data_prompts)

**NO.75** A Salesforce AI Specialist is reviewing the feedback from a customer about the ineffectiveness of the prompt template.

What should the AI Specialist do to ensure the prompt template's effectiveness?

**A.** Monitor and refine the template based on user feedback.

**B.** Use the Prompt Builder Scorecard to help monitor.

**C.** Periodically change the templates grounding object.

**Answer:** B

Explanation:

To address the ineffectiveness of a prompt template reported by a customer, the Salesforce AI Specialist should use the Prompt Builder Scorecard (Option B). This tool is explicitly designed to evaluate and monitor prompt templates against key criteria such as relevance, accuracy, safety, and grounding. By leveraging the scorecard, the specialist can systematically identify weaknesses in the template and make data-driven refinements. While monitoring and refining based on user feedback (Option A) is a general best practice, the Prompt Builder Scorecard is Salesforce's recommended tool for structured evaluation, aligning with documented processes for maintaining prompt effectiveness. Changing the grounding object (Option C) without proper evaluation is reactive and does not address the root cause.

References:

\* Salesforce Einstein AI Specialist Certification Guide: Emphasizes using the Prompt Builder Scorecard to evaluate prompts and iterate based on results.

\* Trailhead Module: "Einstein for Developers" highlights the scorecard as a critical tool for assessing prompt performance.

\* Salesforce Help Documentation: Details the Scorecard's role in evaluating prompts against predefined criteria.

**NO.76** Northern Trail Outfitters (NTO) wants to configure Einstein Trust Layer in its production org but is unable to see the option on the Setup page.

After provisioning Data Cloud, which step must an AI Specialist take to make this option available to NTO?

**A.** Turn on Einstein Copilot.

**B.** Turn on Einstein Generative AI.

**C.** Turn on Prompt Builder.

**Answer:** B

Explanation:

For Northern Trail Outfitters (NTO) to configure the Einstein Trust Layer, the Einstein Generative AI feature must be enabled. The Einstein Trust Layer is closely tied to generative AI capabilities, ensuring that AI-generated content complies with data privacy, security, and trust standards.

\* Option A (Turning on Einstein Copilot) is unrelated to the setup of the Einstein Trust Layer, which focuses more on generative AI interactions and data handling.

\* Option C (Turning on Prompt Builder) is used for configuring and building AI-driven prompts, but it does not enable the Einstein Trust Layer.

Salesforce AI Specialist References: For more details on the Einstein Trust Layer and setup steps:  
[https://help.salesforce.com/s/articleView?id=sf.einstein\\_trust\\_layer\\_overview.htm](https://help.salesforce.com/s/articleView?id=sf.einstein_trust_layer_overview.htm)

**NO.77** Leadership needs to populate a dynamic form field with a summary or description created by a large language model (LLM) to facilitate more productive conversations with customers. Leadership also wants to keep a human in the loop to be considered in their AI strategy.

Which prompt template type should the AI Specialist recommend?

- A. Sales Email
- B. Field Generation
- C. Record Summary

**Answer:** B

Explanation:

The correct answer is Field Generation because this template type is designed to dynamically populate form fields with content generated by a large language model (LLM). In this scenario, leadership wants a dynamic form field that contains a summary or description generated by AI to aid customer interactions. Additionally, they want to keep a human in the loop, meaning the generated content will likely be reviewed or edited by a person before it's finalized, which aligns with the Field Generation prompt template.

\* Field Generation: This prompt type allows you to generate content for specific fields in Salesforce, leveraging large language models to create dynamic and contextual information. It ensures that AI content is available within the record where needed, but it allows human oversight or review, supporting the "human-in-the-loop" strategy.

\* Sales Email: This prompt type is mainly used for generating email content for outreach or responses, which doesn't align directly with populating fields in a form.

\* Record Summary: While this option might seem close, it is typically used to summarize entire records for high-level insights rather than filling specific fields with dynamic content based on AI generation.

Salesforce AI Specialist References:

\* You can explore more about these prompt templates and AI capabilities through Salesforce documentation and official resources on Prompt Builder: [https://help.salesforce.com/s/articleView?id=sf.prompt\\_builder\\_templates\\_overview.htm](https://help.salesforce.com/s/articleView?id=sf.prompt_builder_templates_overview.htm)

**NO.78** Which use case is best supported by Salesforce Einstein Copilot's capabilities?

- A. Bring together a conversational interface for interacting with AI for all Salesforce users, such as developers and ecommerce retailers.
- B. Enable Salesforce admin users to create and train custom large language models (LLMs) using CRM data.
- C. Enable data scientists to train predictive AI models with historical CRM data using built-in machine learning capabilities

**Answer:** A

Explanation:

Salesforce Einstein Copilot is designed to provide a conversational AI interface that can be utilized by different types of Salesforce users, such as developers, sales agents, and retailers. It acts as an AI-powered assistant that facilitates natural interactions with the system, enabling users to perform tasks and access data easily. This includes tasks like pulling reports, updating records, and generating

personalized responses in real time.

\* Option A is correct because Einstein Copilot brings a conversational interface that caters to a wide range of users.

\* Option B and Option C are more focused on developing and training AI models, which are not the primary functions of Einstein Copilot.

References:

\* Salesforce Einstein Copilot Overview: [https://help.salesforce.com/s/articleView?id=einstein\\_copilot\\_overview.htm](https://help.salesforce.com/s/articleView?id=einstein_copilot_overview.htm)

**NO.79** An AI Specialist is tasked with configuring a generative model to create personalized sales emails using customer data stored in Salesforce. The AI Specialist has already fine-tuned a large language model (LLM) on the OpenAI platform. Security and data privacy are critical concerns for the client.

How should the AI Specialist integrate the custom LLM into Salesforce?

**A.** Create an application of the custom LLM and embed it in Sales Cloud via iFrame.

**B.** Add the fine-tuned LLM in Einstein Studio Model Builder.

**C.** Enable model endpoint on OpenAI and make callouts to the model to generate emails.

**Answer:** B

Explanation:

Since security and data privacy are critical, the best option for the AI Specialist is to integrate the fine-tuned LLM (Large Language Model) into Salesforce by adding it to Einstein Studio Model Builder.

Einstein Studio allows organizations to bring their own AI models (BYOM), ensuring the model is securely managed within Salesforce's environment, adhering to data privacy standards.

\* Option A (embedding via iFrame) is less secure and doesn't integrate deeply with Salesforce's data and security models.

\* Option C (making callouts to OpenAI) raises concerns about data privacy, as sensitive Salesforce data would be sent to an external system.

Einstein Studio provides the most secure and seamless way to integrate custom AI models while maintaining control over data privacy and compliance. More details can be found in Salesforce's Einstein Studio documentation on integrating external models.

**NO.80** Where should the AI Specialist go to add/update actions assigned to a copilot?

**A.** Copilot Actions page, the record page for the copilot action, or the Copilot Action Library tab

**B.** Copilot Actions page or Global Actions

**C.** Copilot Detail page, Global Actions, or the record page for the copilot action

**Answer:** A

Explanation:

To add or update actions assigned to a copilot, an AI Specialist can manage this through several areas:

\* Copilot Actions Page: This is the central location where copilot actions are managed and configured

\* Record Page for the Copilot Action: From the record page, individual copilot actions can be updated or modified.

\* Copilot Action Library Tab: This tab serves as a repository where predefined or custom actions for Copilot can be accessed and modified.

These areas provide flexibility in managing and updating the actions assigned to Copilot, ensuring that the AI assistant remains aligned with business requirements and processes.

The other options are incorrect:

- \* B misses the Copilot Action Library, which is crucial for managing actions.
- \* C includes the Copilot Detail page, which isn't the primary place for action management.

References:

- \* Salesforce Documentation on Managing Copilot Actions
- \* Salesforce AI Specialist Guide on Copilot Action Management

**NO.81** An AI Specialist turned on Einstein Generative AI in Setup. Now, the AI Specialist would like to create custom prompt templates in Prompt Builder. However, they cannot access Prompt Builder in the Setup menu.

What is causing the problem?

- A.** The Prompt Template User permission set was not assigned correctly.
- B.** The Prompt Template Manager permission set was not assigned correctly.
- C.** The large language model (LLM) was not configured correctly in Data Cloud.

**Answer:** B

Explanation:

In order to access and create custom prompt templates in Prompt Builder, the AI Specialist must have the Prompt Template Manager permission set assigned. Without this permission, they will not be able to access Prompt Builder in the Setup menu, even though Einstein Generative AI is enabled.

\* Option B is correct because the Prompt Template Manager permission set is required to use Prompt Builder.

\* Option A (Prompt Template User permission set) is incorrect because this permission allows users to use prompts, but not create or manage them.

\* Option C (LLM configuration in Data Cloud) is unrelated to the ability to access Prompt Builder.

References:

- \* Salesforce Prompt Builder Permissions: [https://help.salesforce.com/s/articleView?id=sf.prompt\\_builder\\_permissions.htm](https://help.salesforce.com/s/articleView?id=sf.prompt_builder_permissions.htm)

**NO.82** Universal Containers (UC) wants to enable its sales reps to explore opportunities that are similar to previously won opportunities by entering the utterance, "Show me other opportunities like this one." How should UC achieve this in Einstein Copilot?

- A.** Use the standard Copilot action.
- B.** Create a custom Copilot action calling a flow.
- C.** Create a custom Copilot action calling an Apex class.

**Answer:** A

Explanation:

Universal Containers can achieve the request to explore similar opportunities by using the standard Copilot action. Einstein Copilot has built-in actions to handle natural language queries, such as "Show me other opportunities like this one." The standard action will process the query and return results based on predefined matching criteria like opportunity details and past Closed Won deals.

This approach avoids the need to create custom flows or Apex classes, leveraging out-of-the-box functionality.

For further details, refer to Einstein Copilot for Sales documentation regarding standard actions and

natural language processing.

**NO.83** Universal Containers (UC) needs to improve the agent productivity in replying to customer chats.

Which generative AI feature should help UC address this issue?

- A.** Case Summaries
- B.** Service Replies
- C.** Case Escalation

**Answer:** B

Explanation:

\* Service Replies: This generative AI feature automates and assists in generating accurate, contextual, and efficient replies for customer service agents. It uses past interactions, case data, and the context of the conversation to provide draft responses, thereby enhancing productivity and reducing response times.

\* Case Summaries: Summarizes case information but does not assist directly in replying to customer chats.

\* Case Escalation: Refers to moving cases to higher-level support teams but does not address the need to improve chat response productivity.

Thus, Service Replies is the best feature for this requirement as it directly aligns with improving agent efficiency in replying to chats.

**NO.84** An AI Specialist wants to ground a new prompt template with the User related list.

What should the AI Specialist consider?

- A.** The User related list should have View All access.
- B.** The User related list needs to be included on the record page.
- C.** The User related list is not supported in prompt templates.

**Answer:** C

Explanation:

Salesforce has restrictions on which objects and related lists can be used for grounding prompt templates. This is likely due to security and privacy concerns related to user data.

While it might seem intuitive to use the User related list to provide context to the LLM, Salesforce prevents this to ensure that sensitive user information is not inadvertently exposed or misused. Therefore, the AI Specialist needs to explore alternative ways to incorporate the necessary user information into the prompt template, perhaps by using other related objects or fields that are supported.

**NO.85** Universal Container's internal auditing team asks an AI Specialist to verify that address information is properly masked in the prompt being generated.

How should the AI Specialist verify the privacy of the masked data in the Einstein Trust Layer?

- A.** Enable data encryption on the address field
- B.** Review the platform event logs
- C.** Inspect the AI audit trail

**Answer:** C

Explanation:

The AI audit trail in Salesforce provides a detailed log of AI activities, including the data used, its

handling, and masking procedures applied in the Einstein Trust Layer. It allows the AI Specialist to inspect and verify that sensitive data, such as addresses, is appropriately masked before being used in prompts or outputs.

- \* Enable data encryption on the address field: While encryption ensures data security at rest or in transit, it does not verify masking in AI operations.

- \* Review the platform event logs: Platform event logs capture system events but do not specifically focus on the handling or masking of sensitive data in AI processes.

- \* Inspect the AI audit trail: This is the most relevant option, as it provides visibility into how data is processed and masked in AI activities.

**NO.86** Universal Containers needs a tool that can analyze voice and video call records to provide insights on competitor mentions, coaching opportunities, and other key information. The goal is to enhance the team's performance by identifying areas for improvement and competitive intelligence. Which feature provides insights about competitor mentions and coaching opportunities?

**A.** Call Summaries

**B.** Einstein Sales Insights

**C.** Call Explorer

**Answer:** C

Explanation:

For analyzing voice and video call records to gain insights into competitor mentions, coaching opportunities, and other key information, Call Explorer is the most suitable feature. Call Explorer, a part of Einstein Conversation Insights, enables sales teams to analyze calls, detect patterns, and identify areas where improvements can be made. It uses natural language processing (NLP) to extract insights, including competitor mentions and moments for coaching. These insights are vital for improving sales performance by providing a clear understanding of the interactions during calls.

- \* Call Summaries offer a quick overview of a call but do not delve deep into competitor mentions or coaching insights.

- \* Einstein Sales Insights focuses more on pipeline and forecasting insights rather than call-based analysis.

References:

- \* Salesforce Einstein Conversation Insights Documentation:

[https://help.salesforce.com/s/articleView?](https://help.salesforce.com/s/articleView?id=einstein_conversation_insights.htm)

[id=einstein\\_conversation\\_insights.htm](https://help.salesforce.com/s/articleView?id=einstein_conversation_insights.htm)

**NO.87** What is the primary function of the planner service in the Einstein Copilot system?

**A.** Generating record queries based on conversation history

**B.** Offering real-time language translation during conversations

**C.** Identifying copilot actions to respond to user utterances

**Answer:** C

Explanation:

The primary function of the planner service in the Einstein Copilot system is to identify copilot actions that should be taken in response to user utterances. This service is responsible for analyzing the conversation and determining the appropriate actions (such as querying records, generating a response, or taking another action) that the Einstein Copilot should perform based on user input.

**NO.88** What is an appropriate use case for leveraging Agentforce Sales Agent in a sales context?

- A.** Enable a sales team to use natural language to invoke defined sales tasks grounded in relevant data and be able to ensure company policies are applied. conversationally and in the now or work.
- B.** Enable a sales team by providing them with an interactive step-by-step guide based on business rules to ensure accurate data entry into Salesforce and help close deals faster.
- C.** Instantly review and read incoming messages or emails that are then logged to the correct opportunity, contact, and account records to provide a full view of customer interactions and communications.

**Answer:** A

Explanation:

Agentforce Sales Agent is designed to let sales teams perform tasks via natural language commands, leveraging Salesforce data while adhering to policies. For example, agents can ask the AI to "update the opportunity stage to Closed Won" or "generate a quote," with the system enforcing validations and data security. This use case aligns with Salesforce's vision of conversational AI streamlining workflows without compromising compliance.

\* Step-by-step guides (B) are typically handled by tools like Dynamic Forms or Guided Selling, not Agentforce.

\* Logging messages/emails (C) is managed by Email-to-Case or Service Cloud, not a sales-specific AI agent.

**NO.89** An AI Specialist is setting up a new org and needs to ensure that users can create and execute prompt templates. The AI Specialist is unsure which roles are necessary for these tasks. Which permission sets should the AI Specialist assign to users who need to create and execute prompt templates?

- A.** Prompt Template Manager for creating templates and Data Cloud Admin for executing templates
- B.** Prompt Template Manager for creating templates and Prompt Template User for executing templates
- C.** Data Cloud Admin for creating templates and Prompt Template User for executing templates

**Answer:** B

Explanation:

To effectively manage and use prompt templates, two distinct permission sets are required:

\* Prompt Template Manager: This permission set allows users to create prompt templates. It provides the necessary access to define templates, which can be shared and utilized across the organization.

\* Prompt Template User: This permission set is designed for users who need to execute the templates. It provides the ability to interact with pre-designed prompts and generate outcomes based on these templates.

The Data Cloud Admin permission set is not directly relevant to creating or executing prompt templates but is more focused on managing the Data Cloud.

**NO.90** A support team handles a high volume of chat interactions and needs a solution to provide quick, relevant responses to customer inquiries.

Responses must be grounded in the organization's knowledge base to maintain consistency and accuracy.

Which feature in Einstein for Service should the support team use?

- A. Einstein Service Replies
- B. Einstein Reply Recommendations
- C. Einstein Knowledge Recommendations

**Answer: B**

Explanation:

The support team should use Einstein Reply Recommendations to provide quick, relevant responses to customer inquiries that are grounded in the organization's knowledge base. This feature leverages AI to recommend accurate and consistent replies based on historical interactions and the knowledge stored in the system, ensuring that responses are aligned with organizational standards.

\* Einstein Service Replies (Option A) is focused on generating replies but doesn't have the same emphasis on grounding responses in the knowledge base.

\* Einstein Knowledge Recommendations (Option C) suggests knowledge articles to agents, which is more about assisting the agent in finding relevant articles than providing automated or AI-generated responses to customers.

Salesforce AI Specialist References: For more information on Einstein Reply Recommendations:

[https://help.](https://help.salesforce.com/s/articleView?id=sf.einstein_reply_recommendations_overview.htm)

[salesforce.com/s/articleView?id=sf.einstein\\_reply\\_recommendations\\_overview.htm](https://help.salesforce.com/s/articleView?id=sf.einstein_reply_recommendations_overview.htm)

**NO.91** An AI Specialist built a Field Generation prompt template that worked for many records, but users are reporting random failures with token limit errors.

What is the cause of the random nature of this error?

- A. The number of tokens generated by the dynamic nature of the prompt template will vary by record.
- B. The template type needs to be switched to Flex to accommodate the variable amount of tokens generated by the prompt grounding.
- C. The number of tokens that can be processed by the LLM varies with total user demand.

**Answer: A**

Explanation:

The reason behind the token limit errors lies in the dynamic nature of the prompt template used in Field Generation. In Salesforce's AI generative models, each prompt and its corresponding output are subject to a token limit, which encompasses both the input and output of the large language model (LLM). Since the prompt template dynamically adjusts based on the specific data of each record, the number of tokens varies per record. Some records may generate longer outputs based on their data attributes, pushing the token count beyond the allowable limit for the LLM, resulting in token limit errors.

This behavior explains why users experience random failures—it is dependent on the specific data used in each case. For certain records, the combined input and output may fall within the token limit, while for others, it may exceed it. This variation is intrinsic to how dynamic templates interact with large language models.

Salesforce provides guidance in their documentation, stating that prompt template design should take into account token limits and suggests testing with varied records to avoid such random errors. It does not mention switching to Flex template type as a solution, nor does it suggest that token limits fluctuate with user demand.

Token limits are a constant defined by the model itself, independent of external user load.

References:



- \* Salesforce Developer Documentation on Token Limits for Generative AI Models
- \* Salesforce AI Best Practices on Prompt Design (Trailhead or Salesforce blog resources)

**NO.92** Universal Containers implements three custom actions to get three distinct types of sales summaries for its users. Users are complaining that they are not getting the right summary based on their utterances. What should the AI Specialist investigate as the root cause?

- A.** Review that the custom action is assigned to an Agent.
- B.** Review the action Instructions to ensure they are unique.
- C.** Ensure the input and output types are correctly chosen.

**Answer:** B

Explanation:

The root cause of users receiving incorrect sales summaries lies in non-unique action instructions (Option B). In Einstein Bots, custom actions are triggered based on how well user utterances align with the action instructions defined for each action. If the instructions for the three custom actions overlap or lack specificity, the bot's natural language processing (NLP) cannot reliably distinguish between them, leading to mismatched responses.

Steps to Investigate:

- \* Review Action Instructions: Ensure each custom action has distinct, context-specific instructions.

For example:

- \* Action 1: "Summarize quarterly sales by region."
- \* Action 2: "Generate a product-wise sales breakdown for the current fiscal year."
- \* Action 3: "Provide a comparison of sales performance between online and in-store channels."

Ambiguous or overlapping instructions (e.g., "Get sales summary") cause confusion.

- \* Test Utterance Matching: Use Einstein Bot's training tools to validate if user utterances map to the correct action. Overlap indicates instruction ambiguity.

- \* Refine Instructions: Incorporate keywords or phrases unique to each sales summary type to improve intent detection.

Why Other Options Are Incorrect:

- \* A. Assigning actions to an agent is irrelevant, as custom actions are automated bot components.
- \* C. Input/output types relate to data formatting, not intent routing. While important for execution, they don't resolve utterance mismatches.

References:

- \* Einstein Bot Developer Guide: Stresses the need for unique action instructions to avoid intent conflicts.
- \* Trailhead Module: "Build AI-Powered Bots with Einstein" highlights instruction specificity for accurate action triggering.
- \* Salesforce Help Documentation: Recommends testing and refining action instructions to ensure clarity in utterance mapping.

**NO.93** Universal Containers wants to make a sales proposal and directly use data from multiple unrelated objects (standard and custom) in a prompt template.

What should the AI Specialist recommend?

- A.** Create a Flex template to add resources with standard and custom objects as inputs.
- B.** Create a prompt template passing in a special custom object that connects the records temporarily,

**C.** Create a prompt template-triggered flow to access the data from standard and custom objects.

**Answer:** A

Explanation:

Universal Containers needs to generate a sales proposal using data from multiple unrelated standard and custom objects within a prompt template. The most effective way to achieve this is by using a Flex template.

Flex templates in Salesforce allow AI specialists to create prompt templates that can accept inputs from multiple sources, including various standard and custom objects. This flexibility enables the direct use of data from unrelated objects without the need to create intermediary custom objects or complex flows.

References:

\* Salesforce AI Specialist Documentation - Flex Templates: Explains how Flex templates can be utilized to incorporate data from multiple sources, providing a flexible solution for complex data requirements in prompt templates.

**NO.94** Universal Containers has a custom Agent action calling a flow to retrieve the real-time status of an order from the order fulfillment system.

For the given flow, what should the AI Specialist consider about the running user's data access?

**A.** The flow must have the "with sharing" permission selected in the advanced settings for the permissions, field-level security, and sharing settings to be respected.

**B.** The custom action adheres to the permissions, field-level security, and sharing settings configured in the flow.

. The Agent will always run flows in system mode so the running user's data access will not affect the data returned.

**Answer:** B

Explanation:

When a flow is invoked via a custom Agent action, its data access depends on the flow's runtime configuration, not system mode by default. Salesforce flows can be configured to respect the running user's permissions and sharing settings:

\* If the flow is set to "Run as the User Who Launched the Flow" (enabled in Flow Settings), it adheres to the user's permissions, field-level security (FLS), and sharing rules.

\* Option C is incorrect because flows do not always run in system mode unless explicitly configured to do so.

\* Option A is misleading because "with sharing" is an Apex concept, not a flow setting. Flows use runtime settings like FLS and sharing enforcement.

References:

\* Salesforce Help: Flow Runtime and Security Context

\* Flow Settings: "Run with User Permission and Field-Level Security" ensures data access aligns with the user's permissions.

**NO.95** The marketing team at Universal Containers is looking for a way to personalize emails based on customer behavior, preferences, and purchase history.

Why should the team use Einstein Copilot as the solution?

**A.** To generate relevant content when engaging with each customer

**B.** To analyze past campaign performance

**C.** To send automated emails to all customers

**Answer:** A

Explanation:

Einstein Copilot is designed to assist in generating personalized, AI-driven content based on customer data such as behavior, preferences, and purchase history. For the marketing team at Universal Containers, this is the perfect solution to create dynamic and relevant email content. By leveraging Einstein Copilot, they can ensure that each customer receives tailored communications, improving engagement and conversion rates.

\* Option A is correct as Einstein Copilot helps generate real-time, personalized content based on comprehensive data about the customer.

\* Option B refers more to Einstein Analytics or Marketing Cloud Intelligence, and Option C deals with automation, which isn't the primary focus of Einstein Copilot.

References:

\* Salesforce Einstein Copilot Overview: [https://help.salesforce.com/s/articleView?id=einstein\\_copilot\\_overview.htm](https://help.salesforce.com/s/articleView?id=einstein_copilot_overview.htm)

**NO.96** An AI Specialist has created a copilot custom action using flow as the reference action type. However, it is not delivering the expected results to the conversation preview, and therefore needs troubleshooting.

What should the AI Specialist do to identify the root cause of the problem?

**A.** In Copilot Builder within the Dynamic Panel, turn on dynamic debugging to show the inputs and outputs.

**B.** Copilot Builder within the Dynamic Panel, confirm selected action and observe the values in Input and Output sections.

**C.** In Copilot Builder, verify the utterance entered by the user and review session event logs for debug information.

**Answer:** A

Explanation:

When troubleshooting a copilot custom action using flow as the reference action type, enabling dynamic debugging within Copilot Builder's Dynamic Panel is the most effective way to identify the root cause. By turning on dynamic debugging, the AI Specialist can see detailed logs showing both the inputs and outputs of the flow, which helps identify where the action might be failing or not delivering the expected results.

\* Option B, confirming selected actions and observing the Input and Output sections, is useful for monitoring flow configuration but does not provide the deep diagnostic details available with dynamic debugging.

\* Option C, verifying the user utterance and reviewing session event logs, could provide helpful context, but dynamic debugging is the primary tool for identifying issues with inputs and outputs in real time.

Salesforce AI Specialist References: To explore more about dynamic debugging in Copilot Builder, see: [https://help.salesforce.com/s/articleView?id=sf.copilot\\_custom\\_action\\_debugging.htm](https://help.salesforce.com/s/articleView?id=sf.copilot_custom_action_debugging.htm)

**NO.97** An administrator wants to check the response of the Flex prompt template they've built, but the preview button is greyed out.

What is the reason for this?

- A.** The records related to the prompt have not been selected.
- B.** The prompt has not been saved and activated,
- C.** A merge field has not been inserted in the prompt.

**Answer:** A

Explanation:

When the preview button is greyed out in a Flex prompt template, it is often because the records related to the prompt have not been selected. Flex prompt templates pull data dynamically from Salesforce records, and if there are no records specified for the prompt, it can't be previewed since there is no content to generate based on the template.

\* Option B, not saving or activating the prompt, would not necessarily cause the preview button to be greyed out, but it could prevent proper functionality.

\* Option C, missing a merge field, would cause issues with the output but would not directly grey out the preview button.

Ensuring that the related records are correctly linked is crucial for testing and previewing how the prompt will function in real use cases.

Salesforce AI Specialist References: Refer to the documentation on troubleshooting Flex templates here:

[https://help.salesforce.com/s/articleView?id=sf.flex\\_prompt\\_builder\\_troubleshoot.htm](https://help.salesforce.com/s/articleView?id=sf.flex_prompt_builder_troubleshoot.htm)

**NO.98** Universal Containers (UC) has implemented Generative AI within Salesforce to enable summarization of a custom object called Guest. Users have reported mismatches in the generated information.

In refining its prompt design strategy, which key practices should UC prioritize?

- A.** Enable prompt test mode, allocate different prompt variations to a subset of users for evaluation, and standardize the most effective model based on performance feedback.
- B.** Create concise, clear, and consistent prompt templates with effective grounding, contextual role-playing, clear instructions, and iterative feedback.
- C.** Submit a prompt review case to Salesforce and conduct thorough testing in the playground to refine outputs until they meet user expectations.

**Answer:** B

Explanation:

For Universal Containers (UC) to refine its Generative AI prompt design strategy and improve the accuracy of the generated summaries for the custom object Guest, the best practice is to focus on crafting concise, clear, and consistent prompt templates. This includes:

\* Effective grounding: Ensuring the prompt pulls data from the correct sources.

\* Contextual role-playing: Providing the AI with a clear understanding of its role in generating the summary.

\* Clear instructions: Giving unambiguous directions on what to include in the response.

\* Iterative feedback: Regularly testing and adjusting prompts based on user feedback.

\* Option B is correct because it follows industry best practices for refining prompt design.

\* Option A (prompt test mode) is useful but less relevant for refining prompt design itself.

\* Option C (prompt review case with Salesforce) would be more appropriate for technical issues or complex prompt errors, not general design refinement.

References:

\* Salesforce Prompt Design Best Practices: <https://help.salesforce.com/s/articleView?id=sf>.

[prompt\\_design\\_best\\_practices.htm](#)

**NO.99** What is best practice when refining Einstein Copilot custom action instructions?

- A.** Provide examples of user messages that are expected to trigger the action.
- B.** Use consistent introductory phrases and verbs across multiple action instructions.
- C.** Specify the persona who will request the action.

**Answer:** A

Explanation:

When refining Einstein Copilot custom action instructions, it is considered best practice to provide examples of user messages that are expected to trigger the action. This helps ensure that the custom action understands a variety of user inputs and can effectively respond to the intent behind the messages.

\* Option B (consistent phrases) can improve clarity but does not directly refine the triggering logic.

\* Option C (specifying a persona) is not as crucial as giving examples that illustrate how users will interact with the custom action.

For more details, refer to Salesforce's Einstein Copilot documentation on building and refining custom actions.

**NO.100** Universal Containers (UC) wants to use the Draft with Einstein feature in Sales Cloud to create a personalized introduction email.

After creating a proposed draft email, which predefined adjustment should UC choose to revise the draft with a more casual tone?

- A.** Make Less Formal
- B.** Enhance Friendliness
- C.** Optimize for Clarity

**Answer:** A

Explanation:

When Universal Containers uses the Draft with Einstein feature in Sales Cloud to create a personalized email, the predefined adjustment to Make Less Formal is the correct option to revise the draft with a more casual tone. This option adjusts the wording of the draft to sound less formal, making the communication more approachable while still maintaining professionalism.

\* Enhance Friendliness would make the tone more positive, but not necessarily more casual.

\* Optimize for Clarity focuses on making the draft clearer but doesn't adjust the tone.

For more details, see Salesforce documentation on Einstein-generated email drafts and tone adjustments.

**NO.101** Universal Containers (UC) wants to enable its sales team to get insights into product and competitor names mentioned during calls.

How should UC meet this requirement?

- A.** Enable Einstein Conversation Insights, assign permission sets, define recording managers, and customize insights with up to 50 competitor names.
- B.** Enable Einstein Conversation Insights, connect a recording provider, assign permission sets, and customize insights with up to 25 products.
- C.** Enable Einstein Conversation Insights, enable sales recording, assign permission sets, and customize insights with up to 50 products.

**Answer: C**

Explanation:

To provide the sales team with insights into product and competitor names mentioned during calls, Universal Containers should:

- \* Enable Einstein Conversation Insights: Activates the feature that analyzes call recordings for valuable insights.
- \* Enable Sales Recording: Allows calls to be recorded within Salesforce without needing an external recording provider.
- \* Assign Permission Sets: Grants the necessary permissions to sales team members to access and utilize conversation insights.
- \* Customize Insights: Configure the system to track mentions of up to 50 products and 50 competitors, providing tailored insights relevant to the organization's needs.

Option C accurately reflects these steps. Option A mentions defining recording managers but omits enabling sales recording within Salesforce. Option B suggests connecting a recording provider and limits customization to 25 products, which does not fully meet UC's requirements.

References:

- \* Salesforce AI Specialist Documentation - Setting Up Einstein Conversation Insights: Provides instructions on enabling conversation insights and sales recording.
- \* Salesforce Help - Customizing Conversation Insights: Details how to customize insights with up to 50 products and competitors.
- \* Salesforce AI Specialist Exam Guide: Outlines best practices for implementing AI features like Einstein Conversation Insights in a sales context.

**NO.102** Universal Containers (UC) is implementing Einstein Generative AI to improve customer insights and interactions. UC needs audit and feedback data to be accessible for reporting purposes. What is a consideration for this requirement?

- A.** Storing this data requires Data Cloud to be provisioned.
- B.** Storing this data requires a custom object for data to be configured.
- C.** Storing this data requires Salesforce big objects.

**Answer: A**

Explanation:

When implementing Einstein Generative AI for improved customer insights and interactions, the Data Cloud is a key consideration for storing and managing large-scale audit and feedback data. The Salesforce Data Cloud (formerly known as Customer 360 Audiences) is designed to handle and unify massive datasets from various sources, making it ideal for storing data required for AI-powered insights and reporting. By provisioning Data Cloud, organizations like Universal Containers (UC) can gain real-time access to customer data, making it a central repository for unified reporting across various systems.

- \* Audit and feedback data generated by Einstein Generative AI needs to be stored in a scalable and accessible environment, and the Data Cloud provides this capability, ensuring that data can be easily accessed for reporting, analytics, and further model improvement.
- \* Custom objects or Salesforce Big Objects are not designed for the scale or the specific type of real-time, unified data processing required in such AI-driven interactions. Big Objects are more suited for archival data, whereas Data Cloud ensures more robust processing, segmentation, and analysis capabilities.

References:

\* Salesforce Data Cloud Documentation: <https://www.salesforce.com/products/data-cloud/overview/>

\* Salesforce Einstein AI Overview: <https://www.salesforce.com/products/einstein/overview/>

**NO.103** Universal Containers (UC) is looking to improve its sales team's productivity by providing real-time insights and recommendations during customer interactions.

Why should UC consider using Agentforce Sales Agent?

- A.** To track customer interactions for future analysis
- B.** To automate the entire sales process for maximum efficiency
- C.** To streamline the sales process and increase conversion rates

**Answer:** C

Explanation:

Agentforce Sales Agent provides real-time insights and AI-powered recommendations, which are designed to streamline the sales process and help sales representatives focus on key tasks to increase conversion rates.

It offers features like lead scoring, opportunity prioritization, and proactive recommendations, ensuring that sales teams can interact with customers efficiently and close deals faster.

\* Option A: While tracking customer interactions is beneficial, it is only part of the broader capabilities offered by Agentforce Sales Agent and is not the primary objective for improving real-time productivity.

\* Option B: Agentforce Sales Agent does not automate the entire sales process but provides actionable recommendations to assist the sales team.

\* Option C: This aligns with the tool's core purpose of enhancing productivity and driving sales success.

**NO.104** An AI Specialist is creating a custom action for Agentforce.

Which setting should the AI Specialist test and iterate on to ensure the action performs as expected?

- A.** Action Name
- B.** Action Input
- C.** Action Instructions

**Answer:** C

Explanation:

When creating a custom action for Einstein Bots in Salesforce (including Agentforce), Action Instructions are critical for defining how the bot processes and executes the action. These instructions guide the bot on the logic to follow, such as API calls, data transformations, or conditional steps. Testing and iterating on the instructions ensures the bot understands how to handle dynamic inputs, external integrations, and decision-making.

Salesforce documentation emphasizes that Action Instructions directly impact the bot's ability to execute workflows accurately. For example, poorly defined instructions may lead to incorrect API payloads or failure to parse responses. The Einstein Bot Developer Guide highlights that refining instructions is essential for aligning the bot's behavior with business requirements.

In contrast:

\* Action Name (A) is a static identifier and does not affect functionality.

\* Action Input (B) defines parameters passed to the action but does not dictate execution logic.

Thus, iterating on Action Instructions (C) ensures the action performs as expected.

**NO.105** An AI Specialist needs to create a prompt template to fill a custom field named Latest Opportunities Summary on the Account object with information from the three most recently opened opportunities.

How should the AI Specialist gather the necessary data for the prompt template?

- A.** Create a flow to retrieve the opportunity information.
- B.** Select the Account Opportunity object as a resource when creating the prompt template.
- C.** Select the latest Opportunities related list as a merge field.

**Answer:** A

Explanation:

To gather the necessary data for populating the Latest Opportunities Summary custom field on the Account object with information from the three most recently opened opportunities, the AI Specialist should create a flow. A flow can be configured to query and retrieve the required opportunity records based on criteria such as their open date. Once the flow has gathered the necessary data, it can be used in a prompt template or other automation processes to populate the custom field on the Account record.

\* Option A is correct because creating a flow allows for dynamic data retrieval and control over the logic for selecting the most recent opportunities.

\* Option B and Option C do not provide sufficient control or data retrieval capabilities needed for this scenario.

References:

\* Salesforce Flow Documentation: <https://help.salesforce.com/s/articleView?id=sf.flow.htm>

**NO.106** What is an AI Specialist able to do when the "Enrich event logs with conversation data" setting in Einstein Copilot is enabled?

- A.** View the user click path that led to each copilot action.
- B.** View session data including user Input and copilot responses for sessions over the past 7 days.
- C.** Generate details reports on all Copilot conversations over any time period.

**Answer:** B

Explanation:

When the "Enrich event logs with conversation data" setting is enabled in Einstein Copilot, it allows an AI Specialist or admin to view session data, including both the user input and copilot responses from interactions over the past 7 days. This data is crucial for monitoring how the copilot is being used, analyzing its performance, and improving future interactions based on past inputs.

\* This setting enriches the event logs with detailed conversational data for better insights into the interaction history, helping AI specialists track AI behavior and user engagement.

\* Option A, viewing the user click path, focuses on navigation but is not part of the conversation data enrichment functionality.

\* Option C, generating detailed reports over any time period, is incorrect because this specific feature is limited to data for the past 7 days.

Salesforce AI Specialist References: You can refer to this documentation for further insights:

[https://help.](https://help.salesforce.com/s/articleView?id=sf.einstein_copilot_event_logging.htm)

[salesforce.com/s/articleView?id=sf.einstein\\_copilot\\_event\\_logging.htm](https://help.salesforce.com/s/articleView?id=sf.einstein_copilot_event_logging.htm)

**NO.107** Universal Containers Is Interested In Improving the sales operation efficiency by analyzing their data using AI-powered predictions in Einstein Studio.



Which use case works for this scenario?

- A.** Predict customer sentiment toward a promotion message.
- B.** Predict customer lifetime value of an account.
- C.** Predict most popular products from new product catalog.

**Answer:** B

Explanation:

For improving sales operations efficiency, Einstein Studio is ideal for creating AI-powered models that can predict outcomes based on data. One of the most valuable use cases is predicting customer lifetime value, which helps sales teams focus on high-value accounts and make more informed decisions. Customer lifetime value (CLV) predictions can optimize strategies around customer retention, cross-selling, and long-term engagement.

\* Option B is the correct choice as predicting customer lifetime value is a well-established use case for AI in sales.

\* Option A (customer sentiment) is typically handled through NLP models, while Option C (product popularity) is more of a marketing analysis use case.

References:

\* Salesforce Einstein Studio Use Case Overview: [https://help.salesforce.com/s/articleView?id=sf.einstein\\_studio\\_overview](https://help.salesforce.com/s/articleView?id=sf.einstein_studio_overview)

**NO.108** An AI Specialist is creating a custom action in Einstein Copilot.

Which option is available for the AI Specialist to choose for the custom copilot action?

- A.** Apex trigger
- B.** SOQL
- C.** Flows

**Answer:** C

Explanation:

When creating a custom action in Einstein Copilot, one of the available options is to use Flows. Flows are a powerful automation tool in Salesforce, allowing the AI Specialist to define custom logic and actions within the Copilot system. This makes it easy to extend Copilot's functionality without needing custom code.

While Apex triggers and SOQL are important Salesforce tools, Flows are the recommended method for creating custom actions within Einstein Copilot because they are declarative and highly adaptable. For further guidance, refer to Salesforce Flow documentation and Einstein Copilot customization resources.

**NO.109** The AI Specialist of Northern Trail Outfitters reviewed the organization's data masking settings within the Configure Data Masking menu within Setup. Upon assessing all of the fields, a few additional fields were deemed sensitive and have been masked within Einstein's Trust Layer.

Which steps should the AI Specialist take upon modifying the masked fields?

- A.** Turn off the Einstein Trust Layer and turn it on again.
- B.** Test and confirm that the responses generated from prompts that utilize the data and masked data do not adversely affect the quality of the generated response
- C.** Turn on Einstein Feedback so that end users can report if there are any negative side effects on AI features.

**Answer:** B

**Explanation:**

After modifying masked fields in Einstein's Trust Layer, the next important step is to test and confirm that the responses generated by prompts utilizing the newly masked data still meet quality standards. This ensures that masking sensitive information does not negatively impact the usefulness or accuracy of the AI-generated content. Thorough testing helps identify any issues in prompt performance that could arise due to masking, and adjustments can be made if needed.

\* Option B is correct because testing the effects of masking on AI responses is a critical step in ensuring AI continues to function as expected.

\* Option A (turning off and on the Einstein Trust Layer) is unnecessary after changing the masked fields.

\* Option C (turning on Einstein Feedback) allows for user feedback but is not a direct step following field masking modifications.

**References:**

\* Salesforce Einstein Trust Layer Overview: [https://help.salesforce.com/s/articleView?id=sf.einstein\\_trust\\_layer.htm](https://help.salesforce.com/s/articleView?id=sf.einstein_trust_layer.htm)

**NO.110** A sales rep at Universal Containers is extremely busy and sometimes will have very long sales calls on voice and video calls and might miss key details. They are just starting to adopt new generative AI features.

Which Einstein Generative AI feature should an AI Specialist recommend to help the rep get the details they might have missed during a conversation?

**A.** Call Summary

**B.** Call Explorer

**C.** Sales Summary

**Answer:** A

**Explanation:**

For a sales rep who may miss key details during long sales calls, the AI Specialist should recommend the Call Summary feature. Call Summary uses Einstein Generative AI to automatically generate a concise summary of important points discussed during the call, helping the rep quickly review the key information they might have missed.

\* Call Explorer is designed for manually searching through call data but doesn't summarize.

\* Sales Summary is focused more on summarizing overall sales activity, not call-specific content.

For more details, refer to Salesforce's Call Summary documentation on how AI-generated summaries can improve sales rep productivity.

**NO.111** Universal Containers wants to reduce overall agent handling time minimizing the time spent typing routine answers for common questions in-chat, and reducing the post-chat analysis by suggesting values for case fields.

Which combination of Einstein for Service features enables this effort?

**A.** Einstein Service Replies and Work Summaries

**B.** Einstein Reply Recommendations and Case Summaries

**C.** Einstein Reply Recommendations and Case Classification

**Answer:** C

**Explanation:**

Universal Containers aims to reduce overall agent handling time by minimizing the time agents spend

typing routine answers for common questions during chats and by reducing post-chat analysis through suggesting values for case fields.

To achieve these objectives, the combination of Einstein Reply Recommendations and Case Classification is the most appropriate solution.

1. Einstein Reply Recommendations:

- \* Purpose: Helps agents respond faster during live chats by suggesting the best responses based on historical chat data and common customer inquiries.

- \* Functionality:

- \* Real-Time Suggestions: Provides agents with a list of recommended replies during a chat session, allowing them to quickly select the most appropriate response without typing it out manually.

- \* Customization: Administrators can configure and train the model to ensure the recommendations are relevant and accurate.

- \* Benefit: Significantly reduces the time agents spend typing routine answers, thus improving efficiency and reducing handling time.

2. Case Classification:

- \* Purpose: Automatically suggests or populates values for case fields based on historical data and patterns identified by AI.

- \* Functionality:

- \* Field Predictions: Predicts values for picklist fields, checkbox fields, and more when a new case is created.

- \* Automation: Can be set to auto-populate fields or provide suggestions for agents to approve.

- \* Benefit: Reduces the time agents spend on post-chat analysis and data entry by automating the classification and field population process.

Why Options A and B are Less Suitable:

- \* Option A (Einstein Service Replies and Work Summaries):

- \* Einstein Service Replies: Similar to Reply Recommendations but typically used for email and not live chat.

- \* Work Summaries: Provides summaries of customer interactions but does not assist in field value suggestions.

- \* Option B (Einstein Reply Recommendations and Case Summaries):

- \* Case Summaries: Generates a summary of the case details but does not help in suggesting field values.

References:

- \* Salesforce AI Specialist Documentation - Einstein Reply Recommendations:

- \* Details how Reply Recommendations assist agents in providing quick responses during live chats.

- \* Salesforce AI Specialist Documentation - Einstein Case Classification:

- \* Explains how Case Classification predicts and suggests field values to streamline case management.

- \* Salesforce Trailhead - Optimize Service with AI:

- \* Provides an overview of AI features that enhance service efficiency.

**NO.112** Universal Containers implements Custom Copilot Actions to enhance its customer service operations. The development team needs to understand the core components of a Custom Copilot Action to ensure proper configuration and functionality.

What should the development team review in the Custom Copilot Action configuration to identify one of the core components of a Custom Copilot Action?

**A.** Instructions

**B. Output Types****C. Action Triggers****Answer: A**

Explanation:

\* Instructions: This is a core component of Custom Copilot Actions. Instructions tell the AI model what the action should do and how it should be executed. Clear and concise instructions are crucial for the action to function correctly and provide the expected outcome.

Let's look at why the other options are not the primary core component:

\* Output Types: While important for defining the kind of data the action produces, it's not the core defining element of the action itself.

\* Action Triggers: These determine when the action is initiated, but they don't define the core functionality of the action.

**NO.113** Universal Containers recently launched a pilot program to integrate conversational AI into its CRM business operations with Einstein Copilot.

How should the AI Specialist monitor Copilot's usability and the assignment of actions?

**A. Run a report on the Platform Debug Logs.**

**B. Query the Copilot log data using the metadata API.**

**C. Run Einstein Copilot Analytics.**

**Answer: C**

Explanation:

To monitor Einstein Copilot's usability and the assignment of actions, the AI Specialist should run Einstein Copilot Analytics. This feature provides insights into how often Copilot is used, the types of actions it is handling, and overall user engagement with the system. It's the most effective way to track Copilot's performance and usage patterns.

\* Platform Debug Logs are not relevant for tracking user behavior or the assignment of Copilot actions.

\* Querying the Copilot log data via the Metadata API would not provide the necessary insights in a structured manner.

For more details, refer to Salesforce's Copilot Analytics documentation for tracking AI-driven interactions.

**NO.114** What should an AI Specialist consider when using related list merge fields in a prompt template associated with an Account object in Prompt Builder?

**A. The Activities related list on the Account object is not supported because it is a polymorphic field.**

**B. If person accounts have been enabled, merge fields will not be available for the Account object.**

**C. Prompt generation will yield no response when there is no related list associated with an Account in runtime.**

**Answer: A**

Explanation:

When using related list merge fields in a prompt template associated with the Account object in Prompt Builder, the Activities related list is not supported due to it being a polymorphic field.

Polymorphic fields can reference multiple different types of objects, which makes them incompatible with some merge field operations in prompt generation.

\* Option B is incorrect because person accounts do not limit the availability of merge fields for the

Account object.

\* Option C is irrelevant since even if no related lists are available at runtime, the prompt can still generate based on other available data fields.

For more information, refer to Salesforce documentation on supported fields and limitations in Prompt Builder.

**NO.115** An AI Specialist needs to enable the use of Sales Email prompt templates for the sales team. The AI Specialist has already created the templates in Prompt Builder.

According to best practices, which steps should the AI Specialist take to ensure the sales team can use these templates?

- A.** Assign the Prompt Template User permission set and enable Sales Emails in Setup.
- B.** Assign the Prompt Template Manager permission set and enable Sales Emails in setup.
- C.** Assign the Data Cloud Admin permission set and enable Sales Emails in Setup.

**Answer:** A

Explanation:

To enable Sales Email prompt templates:

\* Permission Set: Assign the Prompt Template User permission set to the sales team to grant access to use pre-built templates.

\* Feature Activation: Enable Sales Emails in Salesforce Setup to activate the integration between prompt templates and email workflows.

\* Option B (Manager permission set): Required for creating/modifying templates, not for usage.

\* Option C (Data Cloud Admin): Unrelated to prompt template access.

References:

\* Salesforce Help: Prompt Template Permissions

\* Specifies that "Prompt Template User" is required to leverage templates in workflows.

\* Sales Email Setup outlines enabling the feature in Setup.

**NO.116** What is the role of the large language model (LLM) in executing an Einstein Copilot Action?

- A.** Find similar requests and provide actions that need to be executed
- B.** Identify the best matching actions and correct order of execution
- C.** Determine a user's access and sort actions by priority to be executed

**Answer:** B

Explanation:

In Einstein Copilot, the role of the Large Language Model (LLM) is to analyze user inputs and identify the best matching actions that need to be executed. It uses natural language understanding to break down the user's request and determine the correct sequence of actions that should be performed. By doing so, the LLM ensures that the tasks and actions executed are contextually relevant and are performed in the proper order. This process provides a seamless, AI-enhanced experience for users by matching their requests to predefined Salesforce actions or flows.

The other options are incorrect because:

A mentions finding similar requests, which is not the primary role of the LLM in this context.

C focuses on access and sorting by priority, which is handled more by security models and governance than by the LLM.

References:

Salesforce Einstein Documentation on Einstein Copilot Actions

## Salesforce AI Documentation on Large Language Models

**NO.117** Universal Containers (UC) wants to enable its sales team with automatic post-call visibility into mention of competitors, products, and other custom phrases.

Which feature should the AI Specialist set up to enable UC's sales team?

- A.** Call Summaries
- B.** Call Explorer
- C.** Call Insights

**Answer:** C

Explanation:

To enable Universal Containers' sales team with automatic post-call visibility into mentions of competitors, products, and custom phrases, the AI Specialist should set up Call Insights. Call Insights analyzes voice and video calls for key phrases, topics, and mentions, providing insights into critical aspects of the conversation. This feature automatically surfaces key details such as competitor mentions, product discussions, and custom phrases specified by the sales team.

\* Call Summaries provide a general overview of the call but do not specifically highlight keywords or topics.

\* Call Explorer is a tool for navigating through call data but does not focus on automatic insights.

For more information, refer to Salesforce's Call Insights documentation regarding the analysis of call content and extracting actionable information.

**NO.118** An AI Specialist implements Einstein Sales Emails for a sales team. The team wants to send personalized follow-up emails to leads based on their interactions and data stored in Salesforce. The AI Specialist needs to configure the system to use the most accurate and up-to-date information for email generation.

Which grounding technique should the AI Specialist use?

- A.** Ground with Apex Merge Fields
- B.** Ground with Record Merge Fields
- C.** Automatic grounding using Draft with Einstein feature

**Answer:** C

Explanation:

For Einstein Sales Emails to generate personalized follow-up emails, it is crucial to ground the email content with the most up-to-date and accurate information. Grounding refers to connecting the AI model with real-time data. The most appropriate technique in this case is Ground with Record Merge Fields. This method ensures that the content in the emails pulls dynamic and accurate data directly from Salesforce records, such as lead or contact information, ensuring the follow-up is relevant and customized based on the specific record.

\* Record Merge Fields ensure the generated emails are highly personalized using data like lead name, company, or other Salesforce fields directly from the records.

\* Apex Merge Fields are typically more suited for advanced, custom logic-driven scenarios but are not the most straightforward for this use case.

\* Automatic grounding using Draft with Einstein is a different feature where Einstein automatically drafts the email, but it does not specifically ground the content with record-specific data like Record Merge Fields.

References:

\* Salesforce Einstein Sales Emails Documentation:

[https://help.salesforce.com/s/articleView?id=release-notes\\_\\_einstein\\_sales\\_emails.htm](https://help.salesforce.com/s/articleView?id=release-notes__einstein_sales_emails.htm)

**NO.119** Universal Containers (UC) has recently received an increased number of support cases. As a result, UC has hired more customer support reps and has started to assign some of the ongoing cases to newer reps.

Which generative AI solution should the new support reps use to understand the details of a case without reading through each case comment?

- A.** Einstein Copilot
- B.** Einstein Sales Summaries
- C.** Einstein Work Summaries

**Answer:** C

Explanation:

New customer support reps at Universal Containers can use Einstein Work Summaries to quickly understand the details of a case without reading through each case comment. Work Summaries leverage generative AI to provide a concise overview of ongoing cases, summarizing all relevant information in an easily digestible format.

\* Einstein Copilot can assist with a variety of tasks but is not specifically designed for summarizing case details.

\* Einstein Sales Summaries are focused on summarizing sales-related activities, which is not applicable for support cases.

For more details, refer to Salesforce documentation on Einstein Work Summaries.

**NO.120** Universal Containers wants to incorporate CRM data as well-formatted JSON in a prompt to a large language model (LLM).

What is an important consideration for this requirement?

- A.** "CRM data to JSON" checkbox must be selected when creating a prompt template.
- B.** Apex code can be used to return a JSON formatted merge field.
- C.** JSON format should be enabled in Prompt Builder Settings.

**Answer:** B

Explanation:

\* Context of the Question

\* Universal Containers (UC) wants to send well-formatted JSON data in a prompt to a large language model (LLM).

\* The question is about an important technical or design consideration for including CRM data as JSON in that prompt.

\* Why Apex Code for JSON Formatting?

\* Apex to Generate JSON: Salesforce does not have a simple "checkbox" or single setting to "convert CRM data to JSON." Typically, to structure data as JSON in a template, you either:

\* Use an Apex class that queries or processes the data, then returns a JSON string.

\* Use a Flow or formula approach (though complex data structures often require Apex).

\* No Built-In "Enable JSON Format in Prompt Builder": Prompt Builder doesn't have a toggle that automatically transforms data into JSON.

\* Conclusion The practical solution to pass CRM data in JSON format to an LLM is to use Apex code (or a specialized Flow approach) to produce a JSON string, which the prompt can then merge and pass

along. Hence, Option B is correct.

Salesforce AI Specialist References & Documents

\* Salesforce Documentation: Working with JSON in Apex Describes how to serialize and deserialize data using Apex for integration or AI prompts.

\* Salesforce AI Specialist Study Guide Emphasizes the need for custom logic (often in Apex) when complex data transformations (like JSON formatting) are required.

**NO.121** How does the Einstein Trust Layer ensure that sensitive data is protected while generating useful and meaningful responses?

**A.** Masked data will be de-masked during response journey.

**B.** Masked data will be de-masked during request journey.

**C.** Responses that do not meet the relevance threshold will be automatically rejected.

**Answer: A**

Explanation:

The Einstein Trust Layer ensures that sensitive data is protected while generating useful and meaningful responses by masking sensitive data before it is sent to the Large Language Model (LLM) and then de- masking it during the response journey.

How It Works:

\* Data Masking in the Request Journey:

\* Sensitive Data Identification: Before sending the prompt to the LLM, the Einstein Trust Layer scans the input for sensitive data, such as personally identifiable information (PII), confidential business information, or any other data deemed sensitive.

\* Masking Sensitive Data: Identified sensitive data is replaced with placeholders or masks. This ensures that the LLM does not receive any raw sensitive information, thereby protecting it from potential exposure.

\* Processing by the LLM:

\* Masked Input: The LLM processes the masked prompt and generates a response based on the masked data.

\* No Exposure of Sensitive Data: Since the LLM never receives the actual sensitive data, there is no risk of it inadvertently including that data in its output.

\* De-masking in the Response Journey:

\* Re-insertion of Sensitive Data: After the LLM generates a response, the Einstein Trust Layer replaces the placeholders in the response with the original sensitive data.

\* Providing Meaningful Responses: This de-masking process ensures that the final response is both meaningful and complete, including the necessary sensitive information where appropriate.

\* Maintaining Data Security: At no point is the sensitive data exposed to the LLM or any unintended recipients, maintaining data security and compliance.

Why Option A is Correct:

\* De-masking During Response Journey: The de-masking process occurs after the LLM has generated its response, ensuring that sensitive data is only reintroduced into the output at the final stage, securely and appropriately.

\* Balancing Security and Utility: This approach allows the system to generate useful and meaningful responses that include necessary sensitive information without compromising data security.

Why Options B and C are Incorrect:

\* Option B (Masked data will be de-masked during request journey):

\* Incorrect Process: De-masking during the request journey would expose sensitive data before it



reaches the LLM, defeating the purpose of masking and compromising data security.

\* Option C (Responses that do not meet the relevance threshold will be automatically rejected):

\* Irrelevant to Data Protection: While the Einstein Trust Layer does enforce relevance thresholds to filter out inappropriate or irrelevant responses, this mechanism does not directly relate to the protection of sensitive data. It addresses response quality rather than data security.

References:

\* Salesforce AI Specialist Documentation - Einstein Trust Layer Overview:

\* Explains how the Trust Layer masks sensitive data in prompts and re-inserts it after LLM processing to protect data privacy.

\* Salesforce Help - Data Masking and De-masking Process:

\* Details the masking of sensitive data before sending to the LLM and the de-masking process during the response journey.

\* Salesforce AI Specialist Exam Guide - Security and Compliance in AI:

\* Outlines the importance of data protection mechanisms like the Einstein Trust Layer in AI implementations.

Conclusion:

The Einstein Trust Layer ensures sensitive data is protected by masking it before sending any prompts to the LLM and then de-masking it during the response journey. This process allows Salesforce to generate useful and meaningful responses that include necessary sensitive information without exposing that data during the AI processing, thereby maintaining data security and compliance.

**NO.122** A data scientist needs to view and manage models in Einstein Studio. The data scientist also needs to create prompt templates in Prompt Builder.

Which permission sets should an AI Specialist assign to the data scientist?

**A.** Data Cloud Admin and Prompt Template Manager

**B.** Prompt Template Manager and Prompt Template User

**C.** Prompt Template User and Data Cloud Admin

**Answer: A**

Explanation:

To allow a data scientist to view and manage models in Einstein Studio and create prompt templates in Prompt Builder, the AI Specialist should assign the Data Cloud Admin and Prompt Template Manager permission sets.

\* Data Cloud Admin provides access to manage and oversee models within Einstein Studio.

\* Prompt Template Manager gives the user the ability to create and manage prompt templates within Prompt Builder.

\* Option A is correct because it assigns the necessary permissions for both managing models and creating prompt templates.

\* Option B and Option C are incorrect as they do not provide the correct combination of permissions for managing models and building prompts.

References:

\* Salesforce Permissions Documentation: [https://help.salesforce.com/s/articleView?id=sf.perm\\_sets\\_overview.htm](https://help.salesforce.com/s/articleView?id=sf.perm_sets_overview.htm)

**NO.123** Universal Containers needs its sales reps to be able to only execute prompt templates.

What should an AI Specialist recommend to achieve this requirement?

**A.** Prompt Template user permission set

**B.** Prompt Template Manager permission set

**C.** Prompt Execute Template permission set

**Answer:** C

Explanation:

\* Prompt Execute Template permission set: This permission set is specifically designed to allow users to execute existing prompt templates. This is exactly what Universal Containers needs for its sales reps.

\* Prompt Template user permission set: This permission set likely grants broader access, including potentially creating or modifying templates, which is not required in this scenario.

\* Prompt Template Manager permission set: This permission set likely grants even more extensive administrative access to prompt templates, going beyond what the sales reps need.

By granting sales reps the "Prompt Execute Template" permission set, you ensure they have the necessary access to use prompt templates without granting unnecessary permissions that could potentially lead to unintended changes or security risks.

**NO.124** In addition to Recipient and Sender, which object should an AI Specialist utilize for inserting merge fields into a Sales email template prompt?

**A.** Recipient Opportunities

**B.** Recipient Account

**C.** User Organization

**Answer:** B

\* Sales Email Template Use Case: When creating a Sales email template (especially for outreach or follow-up), you often need to reference relevant details about the Account linked to the recipient.

\* Standard Merge Fields in Salesforce Email Templates:

\* Recipient (Contact, Lead, or Person receiving the email)

\* Sender (User sending the email)

\* Recipient Account (the Account related to that Contact, providing company-level details and other relevant data)

\* Why Recipient Account?

\* For Sales communications, referencing the Account data (e.g., Account name, industry, or other custom fields) in an email is very common.

\* This is especially important for B2B scenarios where the Contact is tied to an Account.

\* "Recipient Opportunities" could be multiple, so it's less direct for standard email merges. The "User Organization" is more generic internal information, not typically inserted for personalization to the recipient.

\* References and Study Resources:

\* Salesforce Help & Training # Email Templates: Merge Fields

\* Salesforce Trailhead # "Create and Customize Email Templates in Sales Cloud"

\* Salesforce AI Specialist Study Resources (covers recommended best practices for leveraging standard objects like Account in AI-powered or prompt-based communications)

**NO.125** Universal Containers (UC) has a legacy system that needs to integrate with Salesforce. UC wishes to create a digest of account action plans using the generative API feature.

Which API service should UC use to meet this requirement?

**A.** REST API

**B. Metadata API**

**C. SOAP API**

**Answer: A**

Explanation:

To create a digest of account action plans using the generative API feature, Universal Containers should use the REST API. The REST API is ideal for integrating Salesforce with external systems and enabling interaction with Salesforce data, including generative capabilities like creating summaries or digests. It supports modern web standards and is suitable for flexible, lightweight interactions between Salesforce and legacy systems.

\* Metadata API is used for retrieving and deploying metadata, not for data operations like generating summaries.

\* SOAP API is an older API used for integration but is less flexible compared to REST for this specific use case.

For more details, refer to Salesforce REST API documentation regarding using REST for data integration and generating content.

**NO.126** Which mechanism within the Einstein Trust Layer helps to ensure that personal data is handled in compliance with data protection regulations like GDPR?

**A. Toxicity Scoring**

**B. Data Masking**

**C. Prompt Defense**

**Answer: B**

**NO.127** An AI Specialist is creating a custom action for Agentforce.

Which setting should the AI Specialist test and iterate on to ensure the action performs as expected?

**A. Action Input**

**B. Action Name**

**C. Action Instructions**

**Answer: C**

Explanation:

To ensure a custom action in Agentforce performs as expected, the AI Specialist must focus on Action Instructions. Here's why:

\* Action Instructions define the logic, parameters, and steps the AI should follow to execute the action.

They include:

\* How input data is processed.

\* API calls or Apex invocations.

\* Conditional logic (e.g., decision trees). Testing and iterating on these instructions ensures alignment with the intended workflow. For example, incorrect API endpoint references or misconfigured parameters in the instructions will cause failures.

\* Action Input (Option A) refers to the data provided to the action. While validating input formats is important, inputs are static once defined. The primary issue lies in whether the instructions correctly use the inputs.

\* Action Name (Option B) is a descriptive label and does not affect functionality.

Salesforce Documentation Support:

\* Salesforce Einstein Bots & Custom Actions Guide highlights that Action Instructions are where the "core logic" resides, requiring rigorous testing (Source: Einstein Bots Developer Guide).

\* Trailhead Module "Build Custom Actions for Einstein Bots" emphasizes refining instructions to handle edge cases and validate outputs (Source: Trailhead).

By iterating on Action Instructions, the AI Specialist ensures the action's logic, integrations, and error handling are robust.

**NO.128** A Salesforce Administrator wants to generate personalized, targeted emails that incorporate customer interaction data. The admin wants to leverage large language models (LLMs) to write the emails, and wants to reuse templates for different products and customers.

Which solution approach should the admin leverage?

- A.** Use sales Email standard templates
- B.** Create a t field Generation prompt template type
- C.** Create a Sales Email prompt template type.

**Answer:** C

Explanation:

To generate personalized emails using LLMs while reusing templates:

\* Sales Email Prompt Template Type (Option C): Designed specifically for generating dynamic email content by combining LLMs with structured templates. It allows admins to define placeholders (e.g., customer name, product details) and reuse templates across scenarios.

\* Option A: Standard email templates lack LLM integration and dynamic personalization.

\* Option B: "t field Generation" is not a valid Salesforce prompt template type.

References:

\* Salesforce Help: Sales Email Prompt Templates

\* Describes using Sales Email prompt templates to "generate targeted emails using dynamic data and LLMs."

**NO.129** A Salesforce Administrator is exploring the capabilities of Einstein Copilot to enhance user interaction within their organization. They are particularly interested in how Einstein Copilot processes user requests and the mechanism it employs to deliver responses. The administrator is evaluating whether Einstein Copilot directly interfaces with a large language model (LLM) to fetch and display responses to user inquiries, facilitating a broad range of requests from users.

How does Einstein Copilot handle user requests In Salesforce?

- A.** Einstein Copilot will trigger a flow that utilizes a prompt template to generate the message.
- B.** Einstein Copilot will perform an HTTP callout to an LLM provider.
- C.** Einstein Copilot analyzes the user's request and LLM technology is used to generate and display the appropriate response.

**Answer:** C

Explanation:

Einstein Copilot is designed to enhance user interaction within Salesforce by leveraging Large Language Models (LLMs) to process and respond to user inquiries. When a user submits a request, Einstein Copilot analyzes the input using natural language processing techniques. It then utilizes LLM technology to generate an appropriate and contextually relevant response, which is displayed directly to the user within the Salesforce interface.

Option C accurately describes this process. Einstein Copilot does not necessarily trigger a flow

(Option A) or perform an HTTP callout to an LLM provider (Option B) for each user request. Instead, it integrates LLM capabilities to provide immediate and intelligent responses, facilitating a broad range of user requests.

References:

\* Salesforce AI Specialist Documentation - Einstein Copilot Overview: Details how Einstein Copilot employs LLMs to interpret user inputs and generate responses within the Salesforce ecosystem.

\* Salesforce Help - How Einstein Copilot Works: Explains the underlying mechanisms of how Einstein Copilot processes user requests using AI technologies.

**NO.130** An AI Specialist at Universal Containers (UC) Is tasked with creating a new custom prompt template to populate a field with generated output. UC enabled the Einstein Trust Layer to ensure AI Audit data is captured and monitored for adoption and possible enhancements. Which prompt template type should the AI Specialist use and which consideration should they review?

- A.** Flex, and that Dynamic Fields is enabled
- B.** Field Generation, and that Dynamic Fields is enabled
- C.** Field Generation, and that Dynamic Forms is enabled

**Answer:** B

Explanation:

When creating a custom prompt template to populate a field with generated output, the most appropriate template type is Field Generation. This template is specifically designed for generating field-specific outputs using generative AI.

Additionally, the AI Specialist must ensure that Dynamic Fields are enabled. Dynamic Fields allow the system to use real-time data inputs from related records or fields when generating content, ensuring that the AI output is contextually accurate and relevant. This is crucial when populating specific fields with AI-generated content, as it ensures the data source remains dynamic and up-to-date.

The Einstein Trust Layer will track and audit the interactions to ensure the organization can monitor AI adoption and make necessary enhancements based on AI usage patterns.

For further reading, refer to Salesforce's guidelines on Field Generation templates and the Einstein Trust Layer.

**NO.131** Universal Containers wants to utilize Einstein for Sales to help sales reps reach their sales quotas by providing AI-generated plans containing guidance and steps for closing deals.

Which feature should the AI Specialist recommend to the sales team?

- A.** Find Similar Deals
- B.** Create Account Plan
- C.** Create Close Plan

**Answer:** C

Explanation:

The "Create Close Plan" feature is designed to help sales reps by providing AI-generated strategies and steps specifically focused on closing deals. This feature leverages AI to analyze the current state of opportunities and generate a plan that outlines the actions, timelines, and key steps required to move deals toward closure. It aligns directly with the sales team's need to meet quotas by offering actionable insights and structured plans.

\* Find Similar Deals (Option A) helps sales reps discover opportunities similar to their current deals

but doesn't offer a plan for closing.

\* Create Account Plan (Option B) focuses on long-term strategies for managing accounts, which might include customer engagement and retention, but doesn't focus on deal closure.

Salesforce AI Specialist References: For more information on using AI for sales, visit:

[https://help.salesforce.com/s/articleView?id=sf.einstein\\_for\\_sales\\_overview.htm](https://help.salesforce.com/s/articleView?id=sf.einstein_for_sales_overview.htm)

**NO.132** An account manager is preparing for an upcoming customer call and wishes to get a snapshot of key data points from accounts, contacts, leads, and opportunities in Salesforce. Which feature provides this?

**A.** Sales Summaries

**B.** Sales Insight Summary

**C.** Work Summaries

**Answer:** B

Explanation:

Sales Insight Summary aggregates key data points from multiple Salesforce objects (accounts, contacts, leads, opportunities) into a consolidated view, enabling account managers to quickly access relevant information for customer calls.

\* Option A (Sales Summaries): Typically refers to Einstein-generated summaries of specific interactions (e.g., emails, calls), not multi-object snapshots.

\* Option C (Work Summaries): Focuses on summarizing customer service interactions (e.g., chat transcripts), not sales data.

\* Option B (Sales Insight Summary): Directly provides a holistic snapshot of sales-related objects, aligning with the scenario.

References:

\* Salesforce Help: Sales Insight Overview

\* Describes Sales Insight Summary as "a unified view of account, contact, and opportunity data for sales readiness."

**NO.133** Universal Containers (UC) is building a Flex prompt template. UC needs to use data returned by the flow in the prompt template. Which flow element should UC use?

**A.** Add Flex Instructions

**B.** Add Prompt Instructions

**C.** Add Flow Instructions

**Answer:** C

Explanation:

\* Context of the Question

\* Universal Containers (UC) wants to build a Flex prompt template that uses data returned by a Flow.

\* "Flex Prompt Templates" allow admins and AI Specialists to incorporate external or dynamic data into generative AI prompts.

\* Why "Add Flow Instructions" Is Needed

\* Passing Flow Data into Prompt Templates: When configuring the prompt, you must specify how data from the running Flow is passed into the Flex template. The designated element for that is typically "Flow Instructions," which map the Flow outputs to the prompt.

\* Other Options:

\* Add Flex Instructions: Typically controls how the AI responds or structures the output, not how to bring Flow data into the template.

\* Add Prompt Instructions: Usually for static or manual instructions that shape the AI's response, rather than referencing dynamic data from the Flow.

\* Outcome

\* "Add Flow Instructions" ensures the prompt can dynamically use the data that the Flow returns-making Option C correct.

Salesforce AI Specialist References & Documents

\* Salesforce Help & Training: Using Prompt Templates with FlowExplains how to pass Flow variables into a prompt template via a specialized step (e.g., "Flow Instructions").

\* Salesforce AI Specialist Study GuideOutlines how to configure generative AI prompts that reference real-time Flow data.

**NO.134** An AI Specialist at Universal Containers is trying to set up a new Field Generation prompt template. They take the following steps.

1. Create a new Field Generation prompt template.

2. Choose Case as the object type.

3. Select the custom field AI\_Analysis\_c as the target field.

After creating the prompt template, the AI Specialist saves, tests, and activates it. Howsoever, when they go to a case record, the AI Analysis field does not show the (Sparkle) icon on the Edit pencil.

When the AI Specialist was editing the field, it was behaving as a normal field.

Which critical step did the AI Specialist miss?

**A.** They forgot to reactivate the Lightning page layout for the Case object after activating their Field Generation prompt template.

**B.** They forgot that the Case Object is not supported for Add generation as Feinstein Service Replies should be used instead.

**C.** They forgot to edit the Lightning page layout and associate the field to a prompt template

**Answer:** C

Explanation:

For Field Generation prompt templates to display the Sparkle icon (indicating AI-generated content), the target field must be explicitly associated with the prompt template on the Lightning page layout. Even if the prompt template is activated, failing to add the field to the page layout and link it to the template will result in the field behaving as a standard field. Salesforce documentation emphasizes that page layout configuration is mandatory to enable AI-driven field interactions.

\* Reactivating the layout (A) is unnecessary unless the layout itself was modified after activation.

\* Case objects are supported for Field Generation (B is incorrect).

**NO.135** Universal Containers is rolling out a new generative AI initiative.

Which Prompt Builder limitations should the AI Specialist be aware of?

**A.** Rich text area fields are only supported in Flex template types.

**B.** Creations or updates to the prompt templates are not recorded in the Setup Audit Trail.

**C.** Custom objects are supported only for Flex template types.

**Answer:** C

Explanation:

The Prompt Builder in Salesforce has some specific limitations, one of which is that custom objects are supported only for Flex template types. This means that users must rely on Flex templates to integrate custom objects into their prompts.

\* Option A: While rich text area fields have certain restrictions, this does not pertain to the core limitation of integrating custom objects.

\* Option B: Updates and creations for prompt templates are indeed recorded in the Setup Audit Trail, so this statement is incorrect.

\* Option C: This is the correct answer as it reflects a documented limitation of the Prompt Builder.

**NO.136** Universal Containers (UC) wants to enable its sales team to use AI to suggest recommended products from its catalog.

Which type of prompt template should UC use?

**A.** Record summary prompt template

**B.** Email generation prompt template

**C.** Flex prompt template

**Answer:** C

Explanation:

Universal Containers (UC) wants to enable its sales team to leverage AI to recommend products from its catalog. The best option for this use case is a Flex prompt template.

A Flex prompt template is designed to provide flexible, customizable AI-driven recommendations or responses based on specific data points, such as product information, customer needs, or sales history. This template type allows the AI to consider various inputs and parameters, making it ideal for generating product recommendations dynamically.

In contrast:

\* A Record summary prompt template (Option A) is used to summarize data related to a specific record, such as generating a quick summary of a sales opportunity or account, but not for recommending products.

\* An Email generation prompt template (Option B) is tailored for crafting email content and is not suitable for suggesting products based on a catalog.

Given the need for dynamic recommendations that pull from a product catalog and potentially other sales data, the Flex prompt template is the correct approach.

Salesforce References:

\* Salesforce Prompt Templates Overview: <https://help.salesforce.com/s/articleView?id=000391407&type=1>

\* Flex Prompt Template Usage: [https://developer.salesforce.com/docs/atlas.en-us.salesforce\\_ai.meta/salesforce\\_ai/prompt\\_flex\\_template](https://developer.salesforce.com/docs/atlas.en-us.salesforce_ai.meta/salesforce_ai/prompt_flex_template)

**NO.137** Universal Containers (UC) plans to send one of three different emails to its customers based on the customer's lifetime value score and their market segment.

Considering that UC are required to explain why an e-mail was selected, which AI model should UC use to achieve this?

**A.** Predictive model and generative model

**B.** Generative model

**C.** Predictive model



**Answer: C**

Explanation:

Universal Containers should use a Predictive model to decide which of the three emails to send based on the customer's lifetime value score and market segment. Predictive models analyze data to forecast outcomes, and in this case, it would predict the most appropriate email to send based on customer attributes.

Additionally, predictive models can provide explainability to show why a certain email was chosen, which is crucial for UC's requirement to explain the decision-making process.

\* Generative models are typically used for content creation, not decision-making, and thus wouldn't be suitable for this requirement.

\* Predictive models offer the ability to explain why a particular decision was made, which aligns with UC's needs.

Refer to Salesforce's Predictive AI model documentation for more insights on how predictive models are used for segmentation and decision making.

**NO.138** What is a Salesforce AI Specialist able to configure in Data Masking within the Einstein Trust Layer?

**A.** The profiles exempt from masking

**B.** The encryption keys for masking

**C.** The privacy data entities to be masked

**Answer: C**

Explanation:

In the Einstein Trust Layer, the Salesforce AI Specialist can configure privacy data entities to be masked (Option C). This ensures sensitive or personally identifiable information (PII) is obfuscated when processed by AI models.

\* Data Masking Configuration:

\* The AI Specialist defines which fields or data types (e.g., email, phone number, Social Security Number) should be masked. For example, masking the Email field in a prompt response to protect user privacy.

\* This is done through declarative settings in Salesforce, where entities (standard or custom fields) are flagged for masking.

\* Why Other Options Are Incorrect:

\* A. Profiles exempt from masking: Exemptions are typically managed via permissions (e.g., field-level security), not directly within Einstein Trust Layer's Data Masking settings.

\* B. Encryption keys for masking: Encryption is separate from masking. Masking involves obfuscation (e.g., replacing "john@example.com" with "@"), not encryption, which uses keys to secure data.

References:

\* Einstein Trust Layer Documentation: States that Data Masking allows admins to "define which fields should be masked to protect sensitive data."

\* Trailhead Module: "Einstein Trust Layer Basics" explains configuring privacy entities for masking.

\* Salesforce Help Article: "Secure AI with Einstein Trust Layer" details masking configurations for privacy compliance.

**NO.139** Universal Containers is planning a marketing email about products that most closely match a customer's expressed interests.

What should an AI Specialist recommend to generate this email?

- A. Standard email marketing template using Apex or flows for matching interest in products
- B. Custom sales email template which is grounded with interest and product information
- C. Standard email draft with Einstein and choose standard email template

**Answer:** B

Explanation:

To generate an email about products that closely match a customer's expressed interests, an AI Specialist should recommend using a custom sales email template that is grounded with interest and product information. This ensures that the email content is personalized based on the customer's preferences, increasing the relevance of the marketing message.

Using grounding ensures that the generative AI pulls the correct data related to customer interests and product matches, making the email more effective.

For more information, refer to Salesforce documentation on grounding AI-generated content and email personalization strategies.

**NO.140** An AI Specialist at Universal Containers is working on a prompt template to generate personalized emails for product demonstration requests from customers. It is important for the AI-generated email to adhere strictly to the guidelines, using only associated opportunity information, and to encourage the recipient to take the desired action.

How should the AI Specialist include these instructions on a new line in the prompt template?

- A. Surround them with triple quotes (""").
- B. Make sure merged fields are defined.
- C. Use curly brackets {} to encapsulate instructions.

**Answer:** A

Explanation:

In Salesforce prompt templates, instructions that guide how the Large Language Model (LLM) should generate content (in this case, personalized emails) can be included by surrounding the instruction text with triple quotes ("""). This formatting ensures that the LLM adheres to the specific instructions while generating the email content.

The use of triple quotes allows the AI to understand that the enclosed text is a directive for how to approach the task, such as limiting the content to associated opportunity information or encouraging a specific action from the recipient.

Refer to Salesforce Prompt Builder documentation for detailed instructions on how to structure prompts for generative AI.

**NO.141** Universal Containers has a new AI project.

What should an AI Specialist consider when adding a related list on the Account object to be used in the prompt template?

- A. After selecting a related list from the Account, use the field picker to choose merge fields in Prompt Builder.
- B. Prompt Builder must be used to assign the fields from the related list as a JSON format.
- C. The fields for the related list are based on the default page layout of the Account for the current user.

**Answer:** A

Explanation:

\* Context of the Question Universal Containers (UC) wants to include details from a related list on the

Account object in a prompt template. This is typically done via Prompt Builder in Salesforce's generative AI setup.

\* Prompt Builder Behavior

\* Selecting a Related List: Within Prompt Builder, you can navigate to the object (Account) and choose which related list (e.g., Contacts, Opportunities) you want to reference.

\* Field Picker: Once a related list is chosen, Prompt Builder provides a field picker interface, allowing you to select specific fields from that related list. These fields then become available for merge fields or dynamic insertion within your prompt.

\* Why Option A is Correct

\* Direct Alignment with the Standard Process: The recommended approach in Salesforce's documentation is to select a related list and then use the field picker to add the necessary fields into your AI prompt. This ensures the prompt has exactly the data you need from that related list.

\* Why Not Option B (JSON Formatting)

\* No Mandatory JSON Requirement: Although you can structure data as JSON if you desire advanced formatting, Prompt Builder does not require you to manually assign the fields from the related list in JSON. The platform automatically handles how the data is passed along in the background.

\* Why Not Option C (Default Page Layout)

\* Independent of Page Layout: Prompt Builder does not rely strictly on the default page layout for fields. You can configure the fields you want from the related list, independent of how the user's page layout is set up in the UI.

\* Conclusion Since the official Salesforce approach involves selecting a related list and then using the field picker to insert merge fields, Option A is the correct and verified answer.

Salesforce AI Specialist References & Documents

\* Salesforce Official Documentation: Prompt Builder Basics Explains how to reference objects and related lists when building AI prompts.

\* Salesforce Trailhead: Get Started with Prompt Builder Provides hands-on exercises demonstrating how to pick fields from related objects or lists.

\* Salesforce AI Specialist Study Guide Outlines best practices for referencing related records and fields in generative AI prompts.

**NO.142** After creating a foundation model in Einstein Studio, which hyperparameter should an AI Specialist use to adjust the balance between consistency and randomness of a response?

**A.** Presence Penalty

**B.** Variability

**C.** Temperature

**Answer:** C

Explanation:

The Temperature hyperparameter controls the randomness of model outputs:

\* Low Temperature (e.g., 0.2): More deterministic, consistent responses.

\* High Temperature (e.g., 1.0): More creative, varied responses.

\* Presence Penalty (Option A): Discourages repetition of tokens, unrelated to randomness.

\* Variability (Option B): Not a standard hyperparameter in Einstein Studio.

References:

\* Einstein Studio Documentation: Model Hyperparameters

\* Explicitly states "Temperature adjusts the balance between predictable and random outputs."

**NO.143** Universal Containers wants to be able to detect with a high level confidence if content generated by a large language model (LLM) contains toxic language.

Which action should an AI Specialist take in the Trust Layer to confirm toxicity is being appropriately managed?

- A.** Access the Toxicity Detection log in Setup and export all entries where isToxicityDetected is true.
- B.** Create a flow that sends an email to a specified address each time the toxicity score from the response exceeds a predefined threshold.
- C.** Create a Trust Layer audit report within Data Cloud that uses a toxicity detector type filter to display toxic responses and their respective scores.

**Answer: C**

Explanation:

To ensure that content generated by a large language model (LLM) is appropriately screened for toxic language, the AI Specialist should create a Trust Layer audit report within Data Cloud. By using the toxicity detector type filter, the report can display toxic responses along with their respective toxicity scores, allowing Universal Containers to monitor and manage any toxic content generated with a high level of confidence.

\* Option C is correct because it enables visibility into toxic language detection within the Trust Layer and allows for auditing responses for toxicity.

\* Option A suggests checking a toxicity detection log, but Salesforce provides more comprehensive options via the audit report.

\* Option B involves creating a flow, which is unnecessary for toxicity detection monitoring.

References:

\* Salesforce Trust Layer Documentation: [https://help.salesforce.com/s/articleView?id=sf.einstein\\_trust\\_layer\\_audit.htm](https://help.salesforce.com/s/articleView?id=sf.einstein_trust_layer_audit.htm)

**NO.144** Universal Containers implemented Einstein Copilot for its users.

One user complains that Einstein Copilot is not deleting activities from the past 7 days.

What is the reason for this issue?

- A.** Einstein Copilot Delete Record Action permission is not associated to the user.
- B.** Einstein Copilot does not have the permission to delete the user's records.
- C.** Einstein Copilot does not support the Delete Record action.

**Answer: C**

Explanation:

Einstein Copilot currently supports various actions like creating and updating records but does not support the Delete Record action. Therefore, the user's request to delete activities from the past 7 days cannot be fulfilled using Einstein Copilot.

\* Unsupported Action: The inability to delete records is due to the current limitations of Einstein Copilot's supported actions. It is designed to assist with tasks like data retrieval, creation, and updates, but for security and data integrity reasons, it does not facilitate the deletion of records.

\* User Permissions: Even if the user has the necessary permissions to delete records within Salesforce, Einstein Copilot itself does not have the capability to execute delete operations.

References:

\* Salesforce AI Specialist Documentation - Einstein Copilot Supported Actions:

\* Lists the actions that Einstein Copilot can perform, noting the absence of delete operations.

\* Salesforce Help - Limitations of Einstein Copilot:

\* Highlights current limitations, including unsupported actions like deleting records.

**NO.145** What is the importance of Action Instructions when creating a custom Agent action?

- A.** Action Instructions tell the user how to call this action in a conversation
- B.** Action Instructions tell the large language model (LLM) which action to use.
- C.** Action Instructions define the expected user experience of an action.

**Answer:** C

Explanation:

Action Instructions are critical for defining how a custom Agent action should be executed, ensuring alignment with the intended user experience. They provide step-by-step guidance to the bot or LLM on logic, data handling, and integration workflows, directly impacting how users interact with the action. For example, clear instructions prevent errors in API calls or data processing, ensuring seamless interactions.

Salesforce documentation states that poorly defined instructions lead to mismatched expectations, while well-structured instructions ensure the action behaves predictably. This aligns with delivering a consistent user experience.

\* A refers to user invocation, which is handled by dialogue flows, not instructions.

\* B is incorrect because the LLM selects actions based on context/intent, not instructions.

**NO.146** Universal Containers is using Einstein Copilot for Sales to find similar opportunities to help close deals faster.

The team wants to understand the criteria used by the copilot to match opportunities.

What is one criteria that Einstein Copilot for Sales uses to match similar opportunities?

- A.** Matched opportunities are limited to the same account.
- B.** Matched opportunities were created in the last 12 months.
- C.** Matched opportunities have a status of Closed Won from last 12 months.

**Answer:** C

Explanation:

When Einstein Copilot for Sales matches similar opportunities, one of the primary criteria used is whether the opportunities have a status of Closed Won within the last 12 months. This is a key factor in identifying successful patterns that could help close current deals. By focusing on opportunities that have been recently successful, Einstein Copilot can provide relevant insights and suggestions to sales reps to help them close similar deals faster.

For more information, review Salesforce Einstein Copilot documentation related to opportunity matching and sales success patterns.

**NO.147** Universal Containers (UC) needs to save agents time with AI-generated case summaries. UC has implemented the Work Summary feature.

What does Einstein consider when generating a summary?

- A.** Generation is grounded with conversation context, Knowledge articles, and cases.
- B.** Generation is grounded with existing conversation context only.
- C.** Generation is grounded with conversation context and Knowledge articles.

**Answer:** A

Explanation:

When generating a Work Summary, Einstein leverages multiple sources of information to provide a

comprehensive and accurate case summary for agents.

\* Conversation Context:

\* Einstein analyzes the details of the customer interaction, including chat or email threads, to extract relevant information for the summary.

\* Knowledge Articles:

\* It considers linked Knowledge Articles or articles referred to during the case resolution process, ensuring the summary incorporates accurate resolutions or additional resources provided to the customer.

\* Cases:

\* Einstein also examines historical cases and related case records to ground the summary in context from past resolutions or interactions.

\* Option A is correct as it includes all three: conversation context, Knowledge articles, and cases.

\* Option B is incorrect because it limits the grounding to conversation context only, excluding other critical elements.

\* Option C is incorrect because it omits case data, which Einstein considers for more accurate and contextually rich summaries.

**NO.148** Before activating a custom copilot action, an AI Specialist would like is to understand multiple real-world user utterances to ensure the action being selected appropriately. Which tool should the AI Specialist recommend?

**A.** Model Playground

**B.** Einstein Copilot

**C.** Copilot Builder

**Answer:** C

Explanation:

To understand multiple real-world user utterances and ensure the correct action is selected before activating a custom copilot action, the recommended tool is Copilot Builder. This tool allows AI Specialists to design and test conversational actions in response to user inputs, helping ensure the copilot can accurately handle different user queries and phrases. Copilot Builder provides the ability to test, refine, and improve actions based on real-world utterances.

\* Option C is correct as Copilot Builder is designed for configuring and testing conversational actions.

\* Option A (Model Playground) is used for testing models, not user utterances.

\* Option B (Einstein Copilot) refers to the conversational interface but isn't the right tool for designing and testing actions.

References:

\* Salesforce Copilot Builder Overview: [https://help.salesforce.com/s/articleView?id=sf.einstein\\_copilot\\_builder.htm](https://help.salesforce.com/s/articleView?id=sf.einstein_copilot_builder.htm)

**NO.149** Universal Containers (UC) is using standard Service AI Grounding. UC created a custom rich text field to be used with Service AI Grounding.

What should UC consider when using standard Service AI Grounding?

**A.** Service AI Grounding only works with Case and Knowledge objects.

**B.** Service AI Grounding only supports String and Text Area type fields.

**C.** Service AI Grounding visibility works in system mode.

**Answer:** B

**Explanation:**

Service AI Grounding retrieves data from Salesforce objects to ground AI-generated responses. Key considerations:

- \* **Field Types:** Standard Service AI Grounding supports String and Text Area fields. Custom rich text fields (e.g., RichTextArea) are not supported, making Option B correct.
- \* **Objects:** While Service AI Grounding primarily uses Case and Knowledge objects (Option A), the limitation here is the field type, not the object.
- \* **Visibility:** Service AI Grounding respects user permissions and sharing settings unless overridden (Option C is incorrect).

**References:**

- \* Salesforce Help: Service AI Grounding Requirements
- \* Explicitly states support for "Text Area and String fields" only.

**NO.150** Universal Containers has a strict change management process that requires all possible configuration to be completed in a sandbox which will be deployed to production. The AI Specialist is tasked with setting up Work Summaries for Enhanced Messaging. Einstein Generative AI is already enabled in production, and the Einstein Work Summaries permission set is already available in production.

Which other configuration steps should the AI Specialist take in the sandbox that can be deployed to the production org?

- A.** create custom fields to store Issue, Resolution, and Summary; create a Quick Action that updates these fields: add the Wrap Up component to the Messaging Session record page layout: and create Permission Set Assignments for the intended Agents.
- B.** From the Einstein setup menu, select Turn on Einstein: create custom fields to store Issue, Resolution, and Summary: create a Quick Action that updates these fields: and add the wrap up component to the Messaging session record page layout.
- C.** Create custom fields to store issue, Resolution, and Summary; create a Quick Action that updates these fields: and add the Wrap up component to the Messaging session record page layout.

**Answer: C****Explanation:**

- \* **Context of the Question**
- \* Universal Containers (UC) has a strict change management process that requires all possible configuration be completed in a sandbox and deployed to Production.
- \* Einstein Generative AI is already enabled in Production, and the "Einstein Work Summaries" permission set is already available in Production.
- \* The AI Specialist needs to configure Work Summaries for Enhanced Messaging in the sandbox.
- \* **What Can Actually Be Deployed from Sandbox to Production?**
- \* **Custom Fields:** Metadata that is easily created in sandbox and then deployed.
- \* **Quick Actions:** Also metadata-based and can be deployed from sandbox to production.
- \* **Layout Components:** Page layout changes (such as adding the Wrap Up component) can be added to a change set or deployment package.
- \* **Why Option C is Correct**
- \* **No Need to Turn on Einstein in Sandbox for Deployment:** Einstein Generative AI is already enabled in Production; turning it on in the sandbox is typically a manual step if you want to test, but that step itself is not "deployable" in the sense of metadata.
- \* **Permission Set Assignments** (as in Option A) are not deployable metadata. You can deploy the

Permission Set itself but not the specific user assignments. Since the question specifically asks "Which other configuration steps should be taken in the sandbox that can be deployed to the production org?", user assignment is not one of them.

\* Why Not Option A or B?

\* Option A: Mentions creating permission set assignments for agents. This cannot be directly deployed from sandbox to Production, as permission set assignments are user-specific and considered "data," not metadata.

\* Option B: Mentions "Turn on Einstein." But Einstein Generative AI is already enabled in Production. Additionally, "Turning on Einstein" is typically an org-level setting, not a deployable metadata item.

\* Conclusion The main deployable items you can reliably create and test in a sandbox, and then migrate to Production, are:

\* Custom Fields (Issue, Resolution, Summary).

\* A Quick Action that updates those fields.

\* Page Layout Change to include the Wrap Up component.

Therefore, Option C is correct and focuses on actions that are truly deployable as metadata from a sandbox to Production.

Salesforce AI Specialist References & Documents

\* Salesforce Trailhead: Work Summaries with Einstein GPT Provides an overview of how to configure Work Summaries, including the need for custom fields, quick actions, and UI components.

\* Salesforce Documentation: Deploying Metadata Between Orgs Explains what can and cannot be deployed via change sets (e.g., custom fields, page layouts, quick actions vs. user permission set assignments).

\* Salesforce AI Specialist Study Guide Outlines which Einstein Generative AI and Work Summaries configurations are deployable as metadata.

**NO.151** Which feature in the Einstein Trust Layer helps to minimize the risks of jailbreaking and prompt injection attacks?

**A.** Secure Data Retrieval and Grounding

**B.** Data Masking

**C.** Prompt Defense

**Answer:** C

Explanation:

The Einstein Trust Layer is designed to ensure responsible and compliant AI usage. Data Masking (B) is the mechanism that directly addresses compliance with data protection regulations like GDPR by obscuring or anonymizing sensitive personal data (e.g., names, emails, phone numbers) before it is processed by AI models. This prevents unauthorized exposure of personally identifiable information (PII) and ensures adherence to privacy laws.

Salesforce documentation explicitly states that Data Masking is a core component of the Einstein Trust Layer, enabling organizations to meet GDPR requirements by automatically redacting sensitive fields during AI interactions. For example, masked data ensures that PII is not stored or used in AI model training or inference without explicit consent.

In contrast:

\* Toxicity Scoring (A) identifies harmful or inappropriate content in outputs but does not address data privacy.

\* Prompt Defense (C) guards against malicious prompts or injection attacks but focuses on security rather than data protection compliance.



**NO.152** Universal Containers (UC) wants to use Flow to bring data from unified Data Cloud objects to prompt templates.

Which type of flow should UC use?

- A.** Template-triggered prompt flow
- B.** Unified-object linking flow
- C.** Data Cloud-triggered flow

**Answer:** A

Explanation:

In this scenario, Universal Containers wants to bring data from unified Data Cloud objects into prompt templates, and the best way to do that is through a Data Cloud-triggered flow. This type of flow is specifically designed to trigger actions based on data changes within Salesforce Data Cloud objects.

Data Cloud-triggered flows can listen for changes in the unified data model and automatically bring relevant data into the system, making it available for prompt templates. This ensures that the data is both real-time and up-to-date when used in generative AI contexts.

For more detailed guidance, refer to Salesforce documentation on Data Cloud-triggered flows and Data Cloud integrations with generative AI solutions.

**NO.153** An AI Specialist at Universal Containers (UC) is building with no-code tools only. They have many small accounts that are only touched periodically by a specialized sales team, and UC wants to maximize the sales operations team's time. UC wants to help prep the sales team for the calls by summarizing past purchases, interests in products shown by the Contact captured via Data Cloud, and a recap of past email and phone conversations for which there are transcripts.

Which approach should the AI Specialist recommend to achieve this use case?

- A.** Use a prompt template grounded on CRH and Data Cloud data using standard foundation model.
- B.** Fine-Tune the standard foundational model due to the complexity of the data.
- C.** Deploy UC's own custom foundational model on this data first.

**Answer:** A

Explanation:

For no-code implementations, Prompt Builder allows AI Specialists to create prompt templates that dynamically ground responses in Salesforce CRM data (e.g., past purchases) and Data Cloud insights (e.g., product interests) without custom coding. The standard foundation model (e.g., Einstein GPT) can synthesize this data into summaries, leveraging structured and unstructured sources (e.g., email/phone transcripts). Fine-tuning (B) or custom models (C) require code and are unnecessary here, as the use case does not involve unique data patterns requiring model retraining.

**NO.154** Universal Containers (UC) is implementing generative AI and wants to leverage a prompt template to provide responses to customers that gives personalized product recommendations to website visitors based on their browsing history.

Which initial step should UC take to ensure the chatbot can deliver accurate recommendations?

- A.** Design universal product recommendations.
- B.** Write a response script for the chatbot.
- C.** Collect and analyze browsing data.

**Answer:** C

**Explanation:**

To enable personalized product recommendations using generative AI, the foundational step for Universal Containers (UC) is collecting and analyzing browsing data (Option C). Personalized recommendations depend on understanding user behavior, which requires structured data about their browsing history. Without this data, the AI model lacks the context needed to generate relevant suggestions.

- \* Data Collection: UC must first aggregate browsing data (e.g., pages visited, products viewed, session duration) to build a dataset that reflects user preferences.

- \* Data Analysis: Analyzing this data identifies patterns (e.g., frequently viewed categories) that inform how prompts should be structured to retrieve relevant recommendations.

- \* Grounding in Data: Salesforce's Prompt Templates rely on grounding data to generate accurate outputs. Without analyzing browsing data, the prompt template cannot reference meaningful insights for personalization.

Options A and D are incorrect because:

- \* Universal recommendations (A) ignore personalization, which is the core requirement.

- \* Writing a response script (D) addresses chatbot interaction design, not the accuracy of recommendations.

**References:**

- \* Salesforce AI Specialist Certification Guide: Highlights the importance of grounding prompts in relevant data sources to ensure accuracy.

- \* Trailhead Module: "Einstein for Developers" emphasizes data preparation as a prerequisite for effective AI-driven personalization.

- \* Salesforce Help Documentation: Recommends analyzing user behavior data to tailor generative AI outputs in commerce use cases.

**NO.155** An administrator is responsible for ensuring the security and reliability of Universal Containers' (UC) CRM data. UC needs enhanced data protection and up-to-date AI capabilities. UC also needs to include relevant information from a Salesforce record to be merged with the prompt. Which feature in the Einstein Trust Layer best supports UC's need?

**A.** Data masking

**B.** Dynamic grounding with secure data retrieval

**C.** Zero-data retention policy

**Answer:** B

**Explanation:**

Dynamic grounding with secure data retrieval is a key feature in Salesforce's Einstein Trust Layer, which provides enhanced data protection and ensures that AI-generated outputs are both accurate and securely sourced. This feature allows relevant Salesforce data to be merged into the AI-generated responses, ensuring that the AI outputs are contextually aware and aligned with real-time CRM data.

Dynamic grounding means that AI models are dynamically retrieving relevant information from Salesforce records (such as customer records, case data, or custom object data) in a secure manner. This ensures that any sensitive data is protected during AI processing and that the AI model's outputs are trustworthy and reliable for business use.

The other options are less aligned with the requirement:

- \* Data masking refers to obscuring sensitive data for privacy purposes and is not related to merging Salesforce records into prompts.

\* Zero-data retention policy ensures that AI processes do not store any user data after processing, but this does not address the need to merge Salesforce record information into a prompt.

References:

- \* Salesforce Developer Documentation on Einstein Trust Layer
- \* Salesforce Security Documentation for AI and Data Privacy

**NO.156** Universal Containers' sales team engages in numerous video sales calls with prospects across the nation. Sales management wants an easy way to understand key information such as deal terms or customer sentiments.

Which Einstein Generative AI feature should an AI Specialist recommend for this request?

- A.** Einstein Call Summaries
- B.** Einstein Conversation Insights
- C.** Einstein Video KPI

**Answer:** A

Explanation:

Einstein Call Summaries is the best option for this scenario because it leverages Salesforce's AI capabilities to automatically summarize key details of video or voice calls. It includes details like deal terms, customer sentiments, follow-up tasks, and other crucial information. This feature is designed to help sales teams focus on their strategies rather than taking extensive manual notes during conversations.

- \* Einstein Call Summaries: Automatically generates summaries for calls, identifying critical points such as next steps and follow-ups, enhancing efficiency and understanding of deal progression.
- \* Einstein Conversation Insights: While it provides insights into customer sentiment and engagement, it is more suited for analyzing patterns across conversations rather than summarizing specific call details.
- \* Einstein Video KPI: Focuses on analyzing key performance indicators within video calls but does not offer summarization features needed for deal terms or sentiment tracking. This feature ensures actionable insights are delivered directly into the Salesforce CRM, allowing sales managers to gain a concise overview without manually reviewing long recordings.

**NO.157** An AI Specialist wants to troubleshoot their Agent's performance.

Where should the AI Specialist go to access all user interactions with the Agent, including Agent errors, incorrectly triggered actions, and incomplete plans?

- A.** Event Logs
- B.** Plan Canvas
- C.** Agent Settings

**Answer:** B

Explanation:

Event Logs in Salesforce capture detailed interaction data, including agent errors, triggered actions, and incomplete plans. These logs provide visibility into user-Agent interactions for troubleshooting performance issues. The Einstein Bot Analytics documentation highlights Event Logs as the primary source for auditing bot behavior and diagnosing issues like misconfigured actions or plan execution failures.

- \* Plan Canvas (B) is for designing workflows, not auditing.
- \* Agent Settings (C) control configuration but do not store interaction history.

TOPEXAM.JP